



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# Introduction and Deep Dive Into Containerd

*Kohei Tokunaga & Akihiro Suda, NTT Corporation*





KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# Introduction to containerd

Kohei Tokunaga, NTT Corporation



<https://github.com/containerd/containerd>

- CNCF graduated container runtime project
- Resource manager
  - Container process
  - Image artifacts
  - Filesystem snapshots
  - Metadata and dependencies management
- Tightly scoped (100% approval is required to stretch) but highly extensible
- Used by Kubernetes, Docker and various container-based projects

# Usage in community



KubeCon




CloudNativeCon

Europe 2021

*Virtual*

- Docker's use of containerd + pure use of containerd is 83% of container usage (Sysdig 2021 container security and usage report)
- Used by several managed services as well as open source projects in community



<https://sysdig.com/blog/sysdig-2021-container-security-usage-report/>

## Adoption

- Managed: GKE, AWS Fargate, AKS, IKS
- Development: Docker/moby, BuildKit
- K8s distribution: k3s, kind, minikube, kubespray, microk8s, k0s
- FaaS: faasd

# How containerd is used?



KubeCon

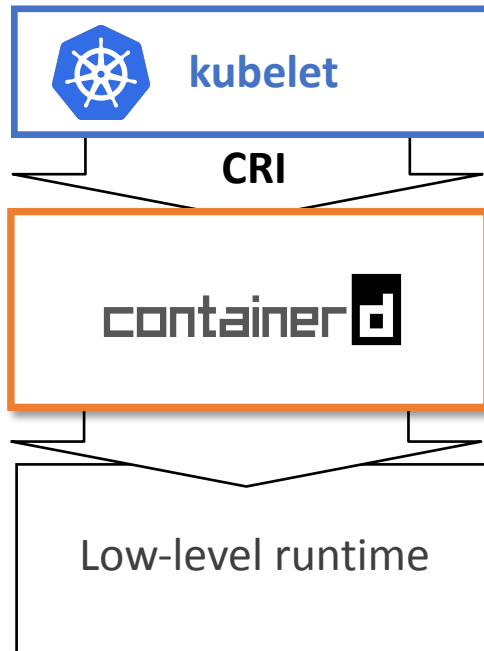


CloudNativeCon

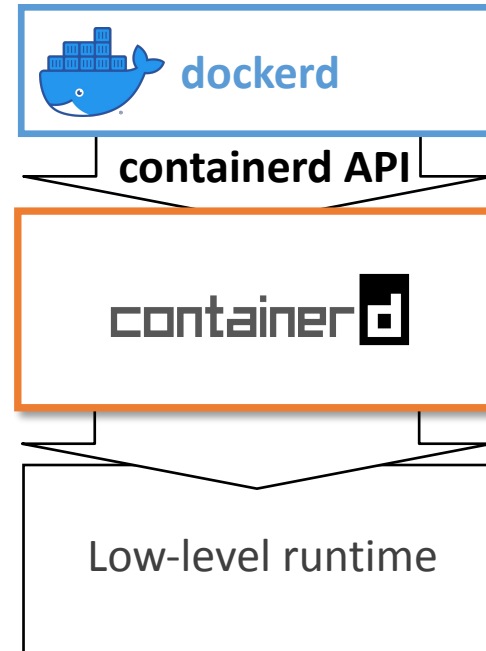
Europe 2021

*Virtual*


As a CRI runtime



As a component of Docker



As a general container management tool



# Containerd as a CRI runtime



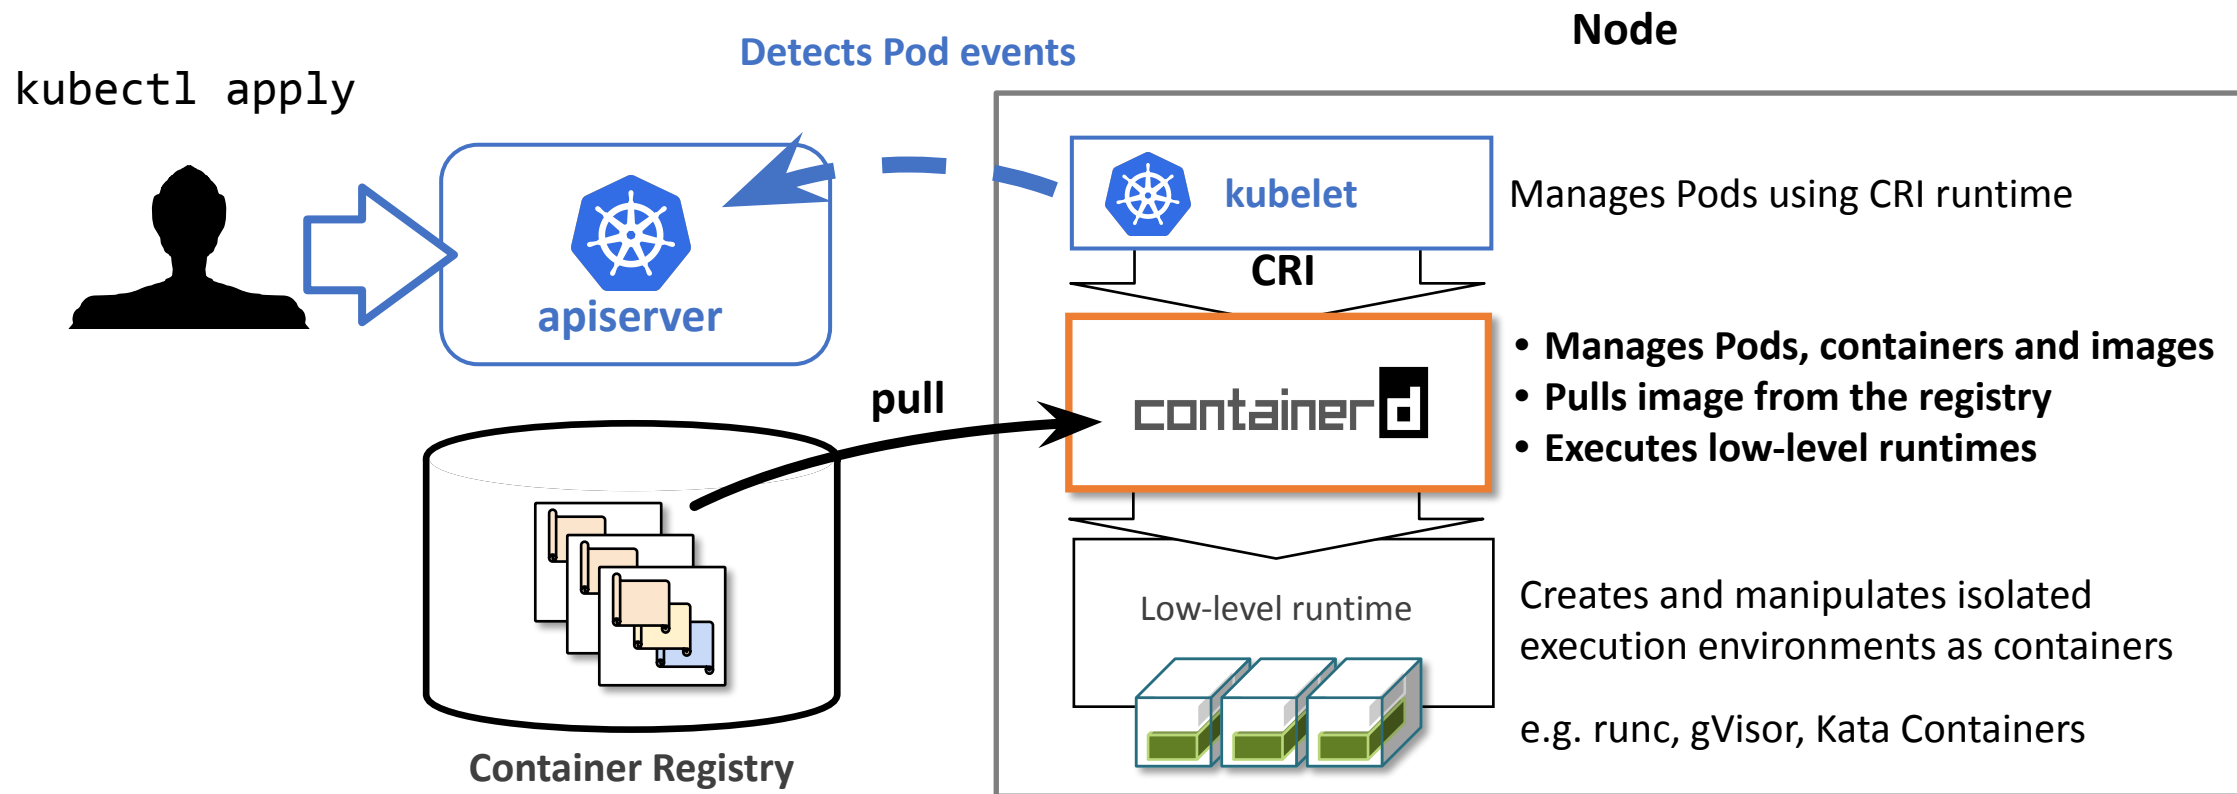
KubeCon



CloudNativeCon

Europe 2021

*Virtual*



The de facto standard CRI runtime for Kubernetes

- Managed Kubernetes: IKS, GKE, AKS, AWS Fargate, ...
- Kubernetes distributions: K3s, kind, minikube, kubespray, microk8s, k0s, ...

# Containerd as a component of Docker



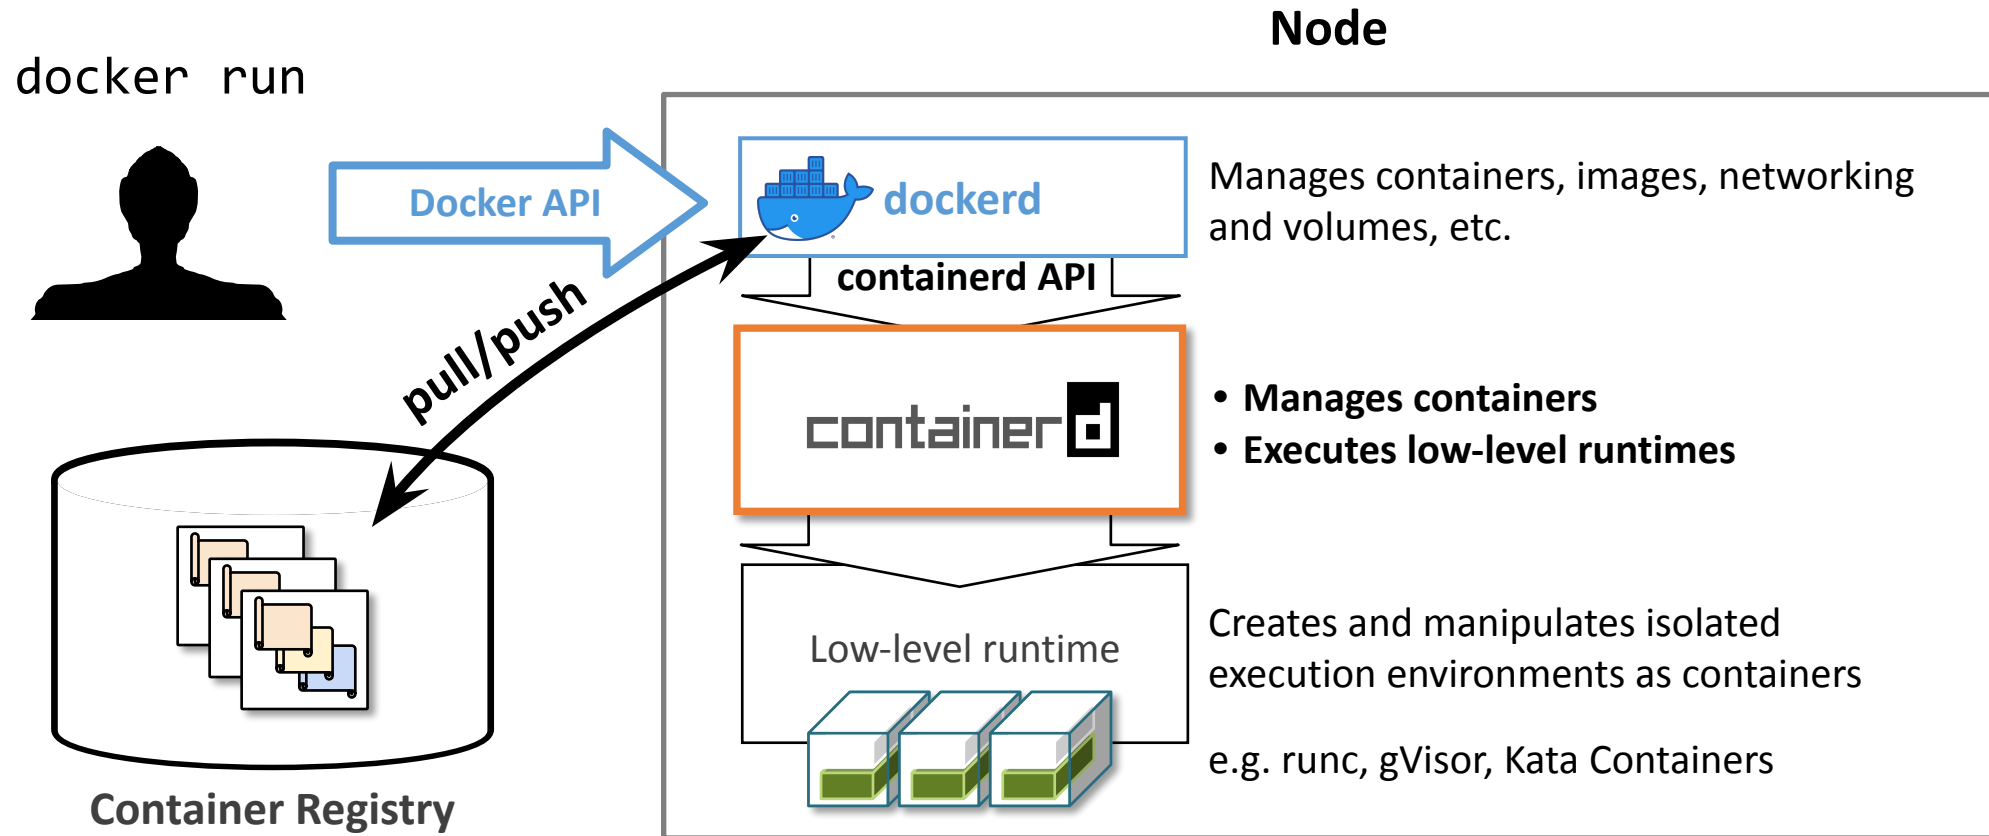
KubeCon



CloudNativeCon

Europe 2021

*Virtual*



# Containerd as a general container management tool



KubeCon

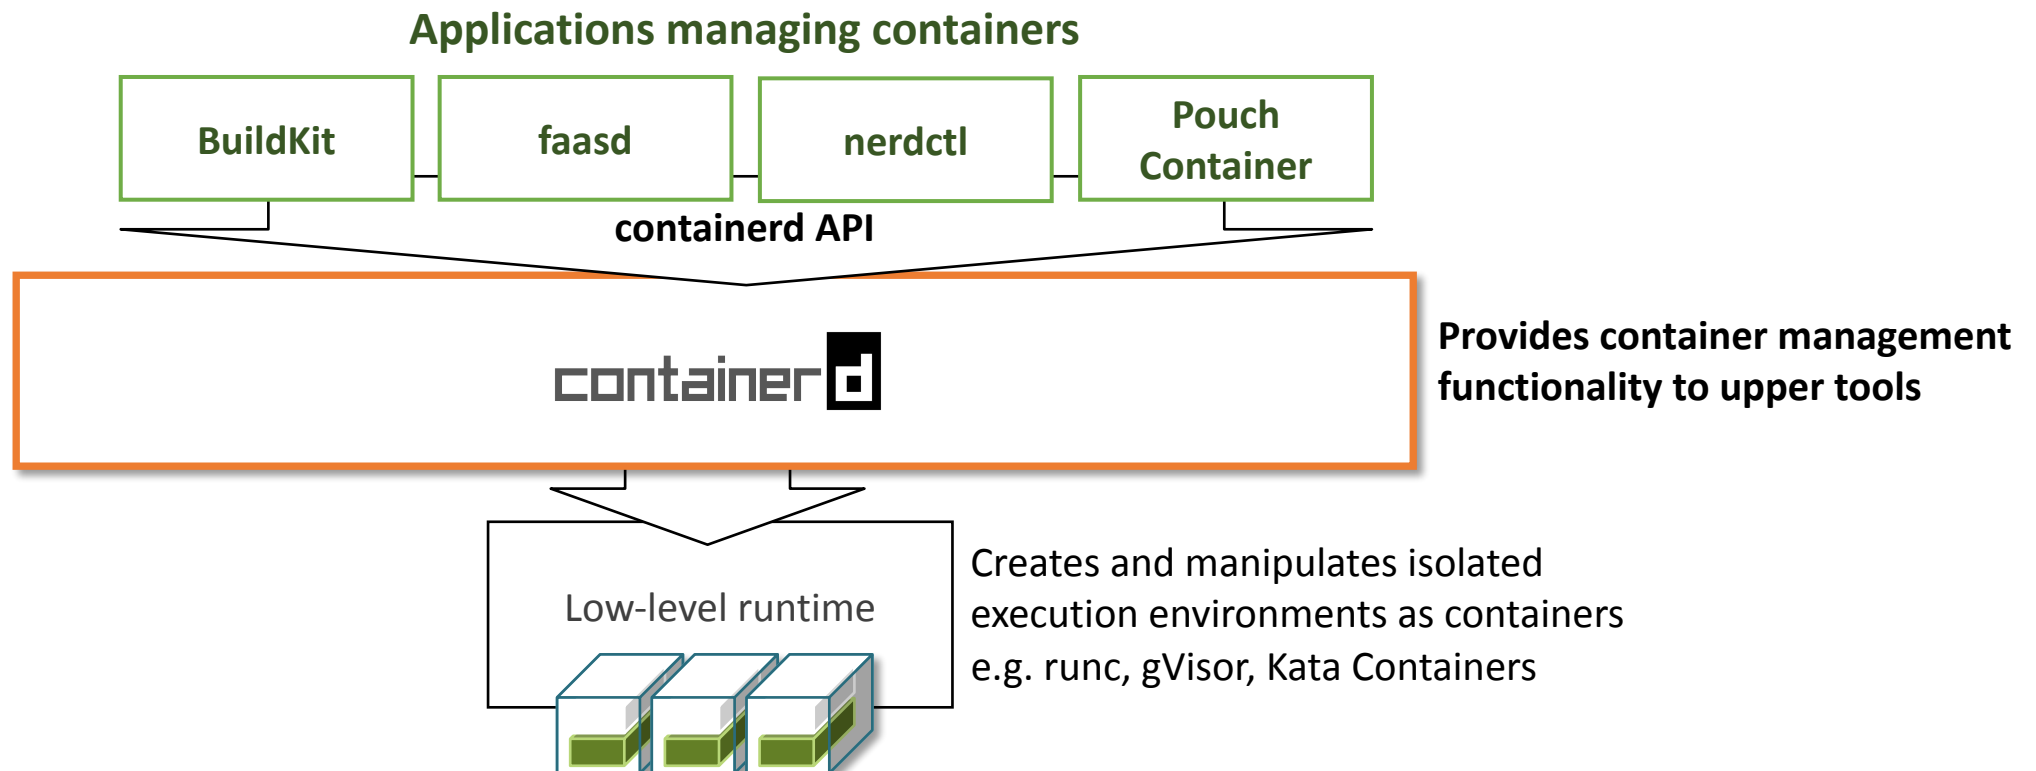


CloudNativeCon

Europe 2021

Virtual

- Several applications are developed based on containerd
- Containerd provides a Go client library (discussed later)
- Applications can extend containerd with plugins, *without recompilation* (discussed later)







KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# Containerd Internal

Kohei Tokunaga, NTT Corporation

# Containerd Architecture



KubeCon

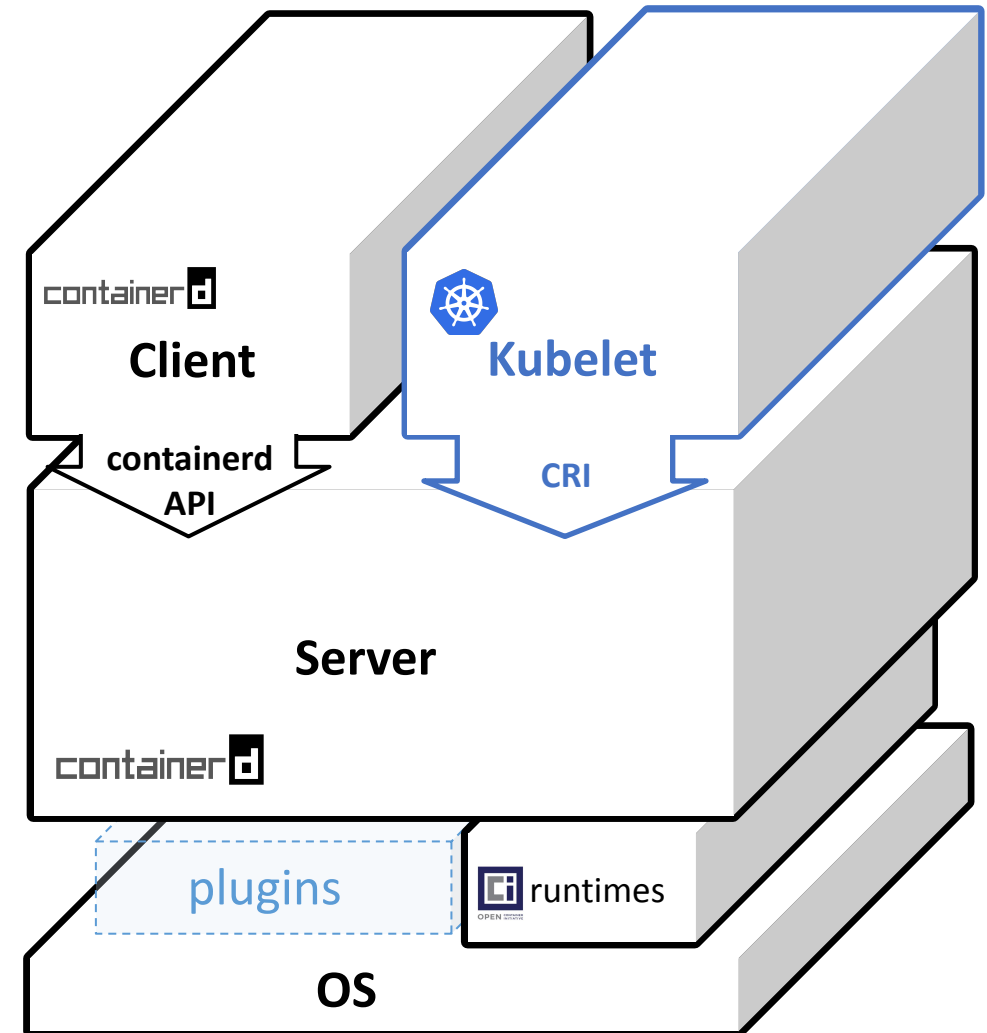


CloudNativeCon

Europe 2021

*Virtual*

- Client-server architecture
  - Go client library (used by Docker, BuildKit, etc.)
- Client calls server via **containerd API**
  - Through `/run/containerd/containerd.sock`
- Various low-level runtimes are supported
  - OCI runtimes (runc, gVisor, Kata Container, etc)
  - Firecracker (firecracker-containerd)
- Extensibility
  - Low-level plugins
  - Extending containerd API with custom services
  - Client library is easy to customize



# Containerd Client



KubeCon

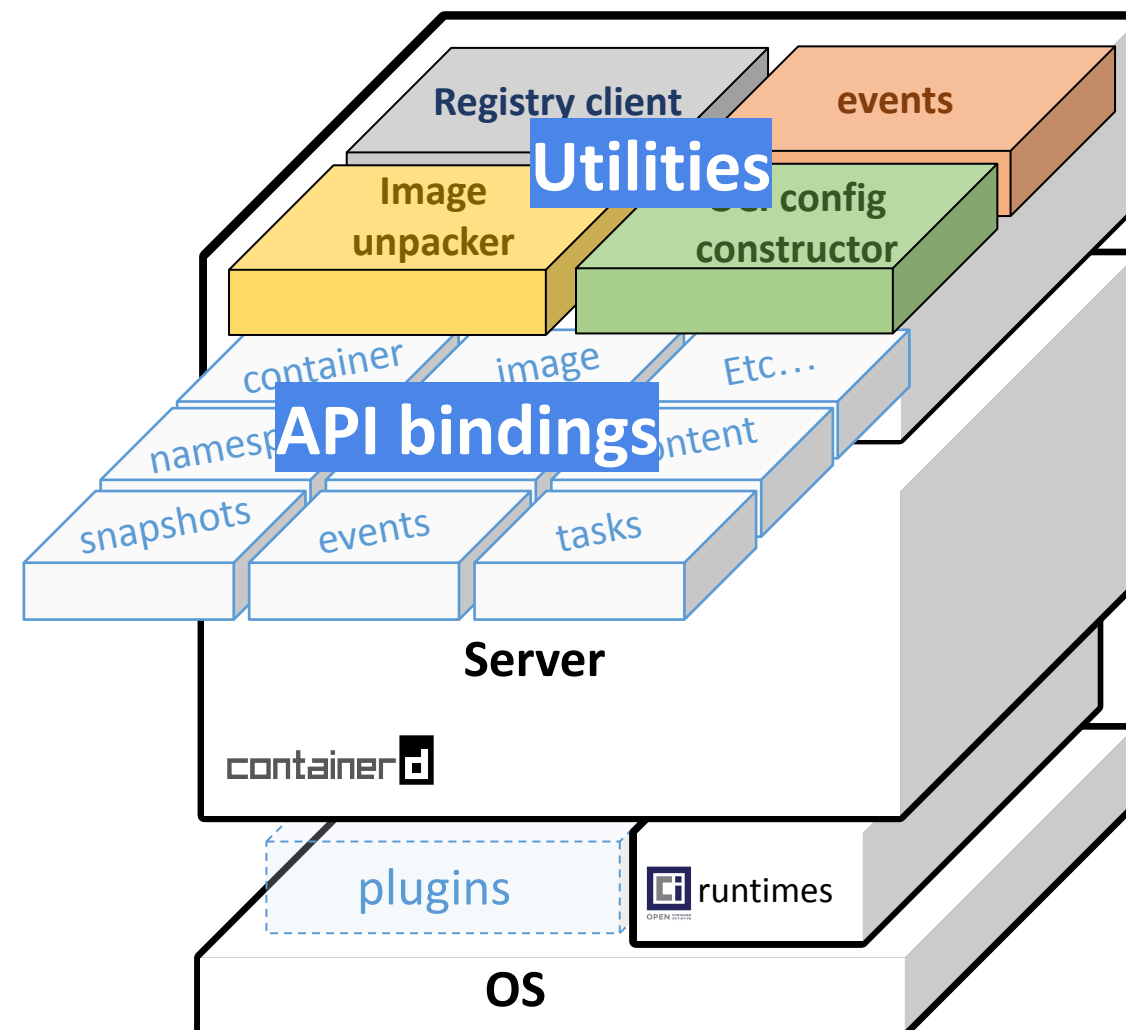


CloudNativeCon

Europe 2021

*Virtual*

- “Smart” Client (Go library)
  - Containerd API bindings
  - Registry client
  - Pulling/Pushing images
  - Image unpacker
  - Creating OCI config for OCI runtimes
- Go application can integrate with containerd using client library



# Containerd Client Implementations



KubeCon




CloudNativeCon

Europe 2021

*Virtual*

- **ctr:** <https://github.com/containerd/containerd>
  - CLI client for containerd
  - Mainly for debugging or trying new features
- **nerdctl:** <https://github.com/containerd/nerdctl>
  - Docker-compatible CLI for containerd
  - Easy to use for Docker users
  - Supports containerd's cutting-edge features (e.g. lazy pulling, image encryption)
- **containerd-based tools**
  - Arbitrary tools can integrate to containerd using client library
  - e.g. Docker, BuildKit, faasd



# Containerd Core & API



KubeCon

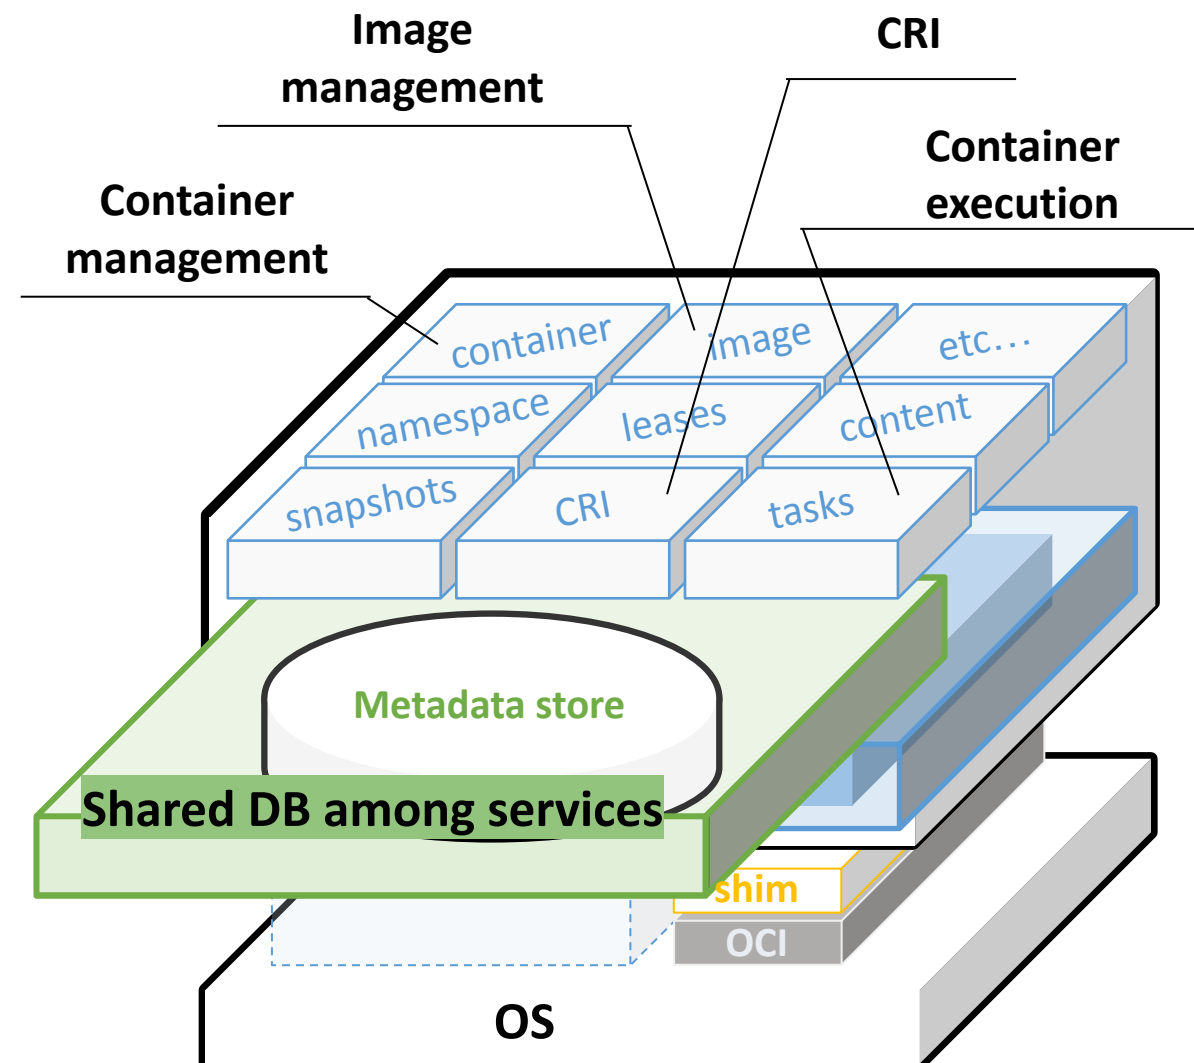


CloudNativeCon

Europe 2021

*Virtual*

- Micro services
  - Containerd API is the set of APIs of services
  - Services are loosely connected
- Shared metadata DB
  - bbolt-based
    - <https://github.com/etcd-io/bbolt>
  - Stores metadata of containers, images, contents, snapshots, etc.
  - Manages reference graph for GC



# CRI Service



KubeCon

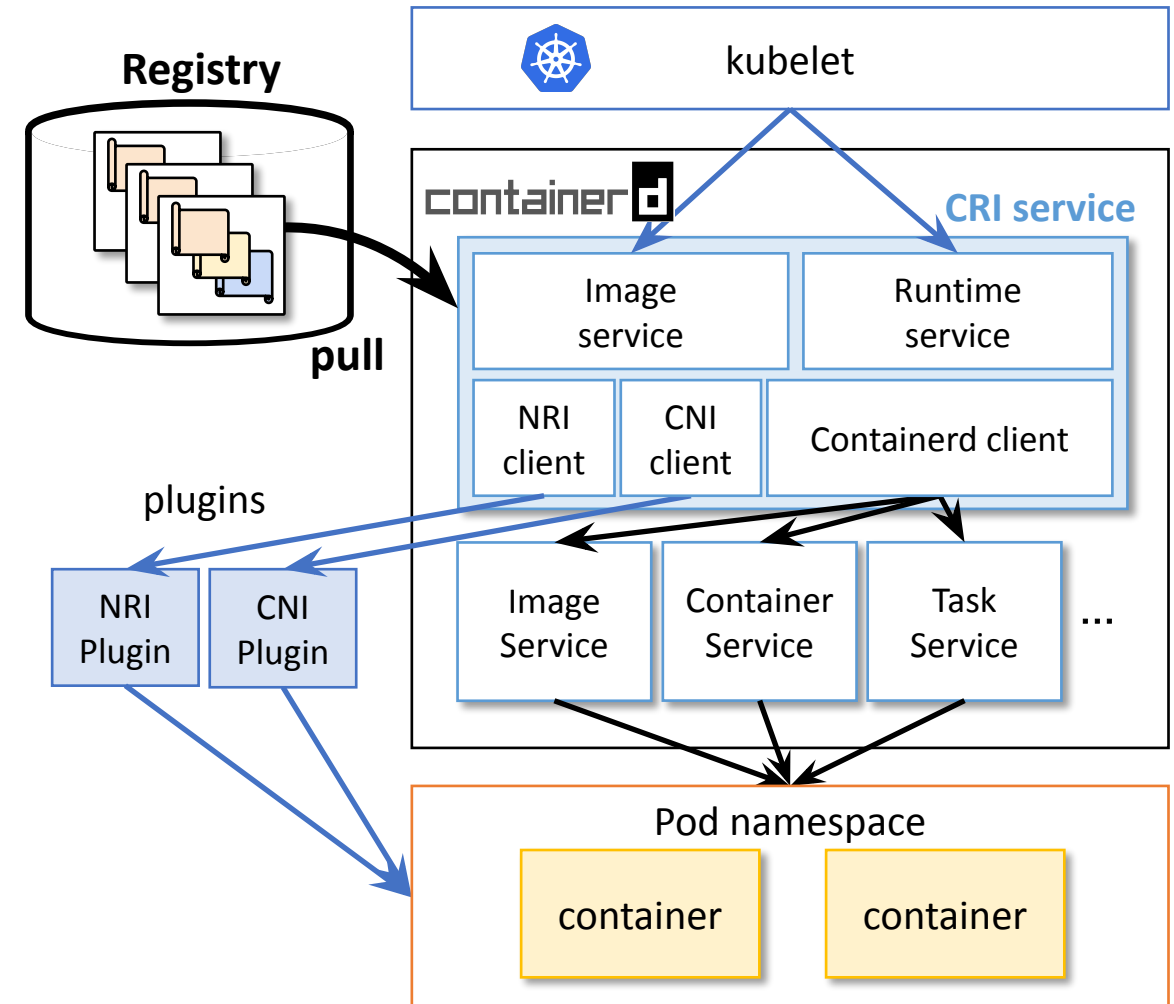


CloudNativeCon

Europe 2021

*Virtual*

- CRI service implements CRI of Kubernetes
- Implemented as a builtin service
  - Initially repo and binary were separated from containerd
  - Merged to containerd/containerd since 1.5
- Depends on other services for container & image management
  - Communicates via function call
- Uses external CNI/NRI plugins for networking and resource management



# Low-level Services



KubeCon




CloudNativeCon

Europe 2021

*Virtual*

- **Content Store**
  - Stores image manifest and layers “as-is”
  - content addressable (keyed by digest)
- **Snapshotter**
  - Manages “snapshots”
    - Extracted and stacked view of rootfs layers
  - Passed to OCI runtimes as rootfs
  - Snapshotter impl. per backing filesystem
    - Overlayfs, btrfs, aufs, FUSE, ...
- **Runtime**
  - Executes low-level runtimes via “shim”
  - Shim is a wrapper daemon of OCI runtime
  - Well-suited to stateful runtimes (e.g. Kata Containers)



# Image content flow



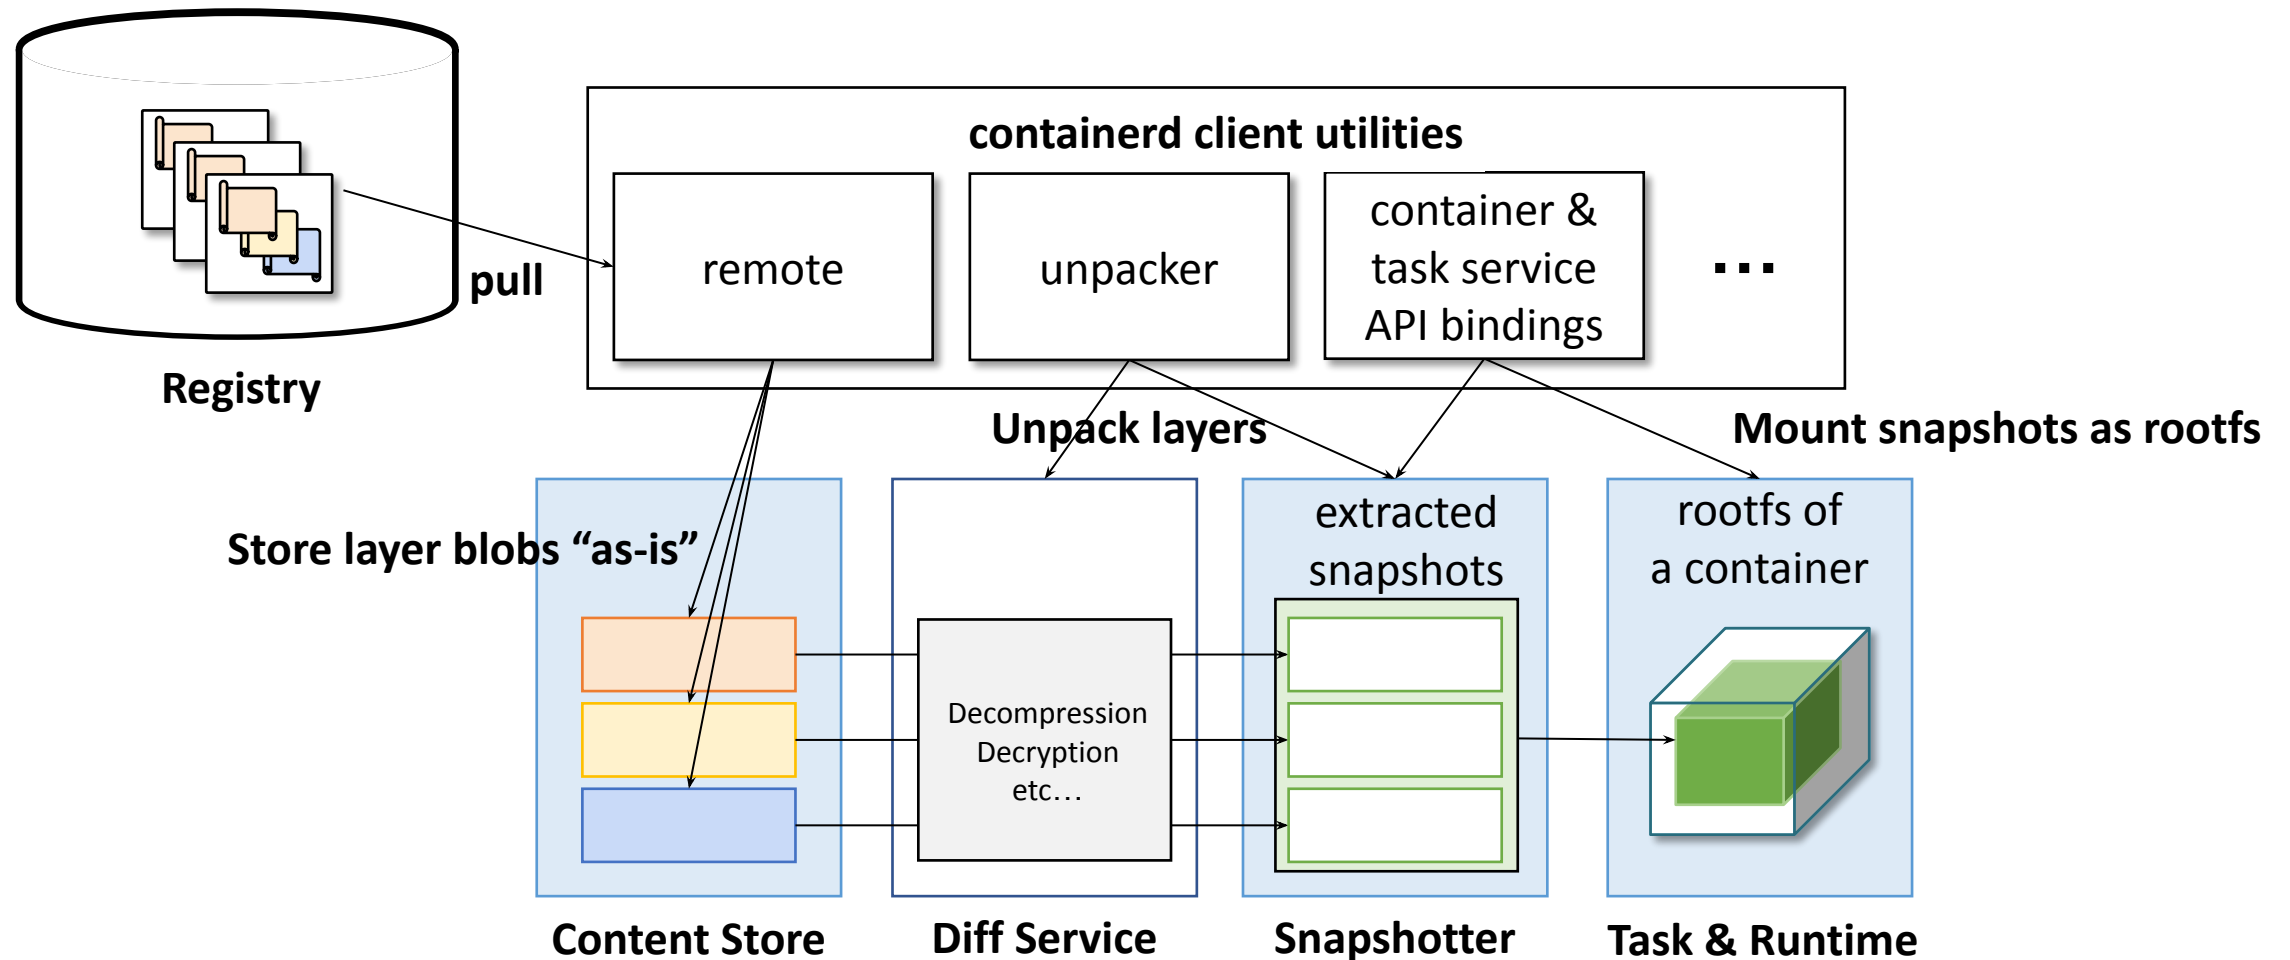
KubeCon



CloudNativeCon

Europe 2021

*Virtual*







KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# Containerd Extensibility

Kohei Tokunaga, NTT Corporation

# Extending containerd with plugins and services



KubeCon




CloudNativeCon

Europe 2021

*Virtual*

- containerd is tightly scoped but highly extensible
- Custom low-level service; **no need to recompile**
  - external binary plugins
    - Plugin via unix socket (proxy snapshotter, proxy content store)
    - Plugin as an executable binary (stream processor, shim)
  - Go plugin
- API is extendable by implementing your own custom service
  - e.g. "control API" of firecracker-containerd



# Extension example 1: Lazy pulling



KubeCon

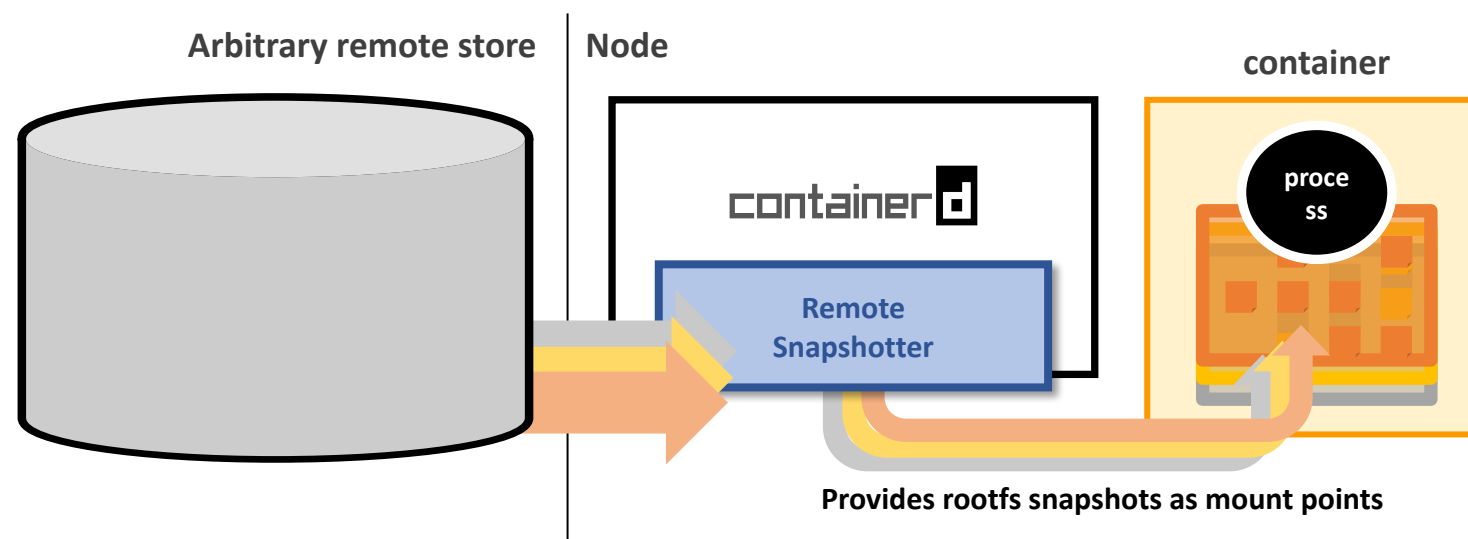


CloudNativeCon

Europe 2021

*Virtual*

- **Remote snapshotter plugin**
  - allows “lazy pulling” of images from arbitrary remote store (not limited to the registry)
    - container can startup without waiting for the entire image contents being locally available
- Snapshotter can run as an external daemon (proxy snapshotter)
  - No re-compilation is required
  - Containerd talks with the snapshotter via unix socket
- **Stargz Snapshotter** enables lazy pulling of OCI-compatible eStargz/Stargz images from standard registry
  - <https://github.com/containerd/stargz-snapshotter>



## Remote Snapshotters in community

- [Stargz Snapshotter](#)
- [CVMFS-snapshotter](#)
- [Nydus-snapshotter](#)
- [OverlayBD-snapshotter](#)

# Extension example 2: Generic image layers



KubeCon

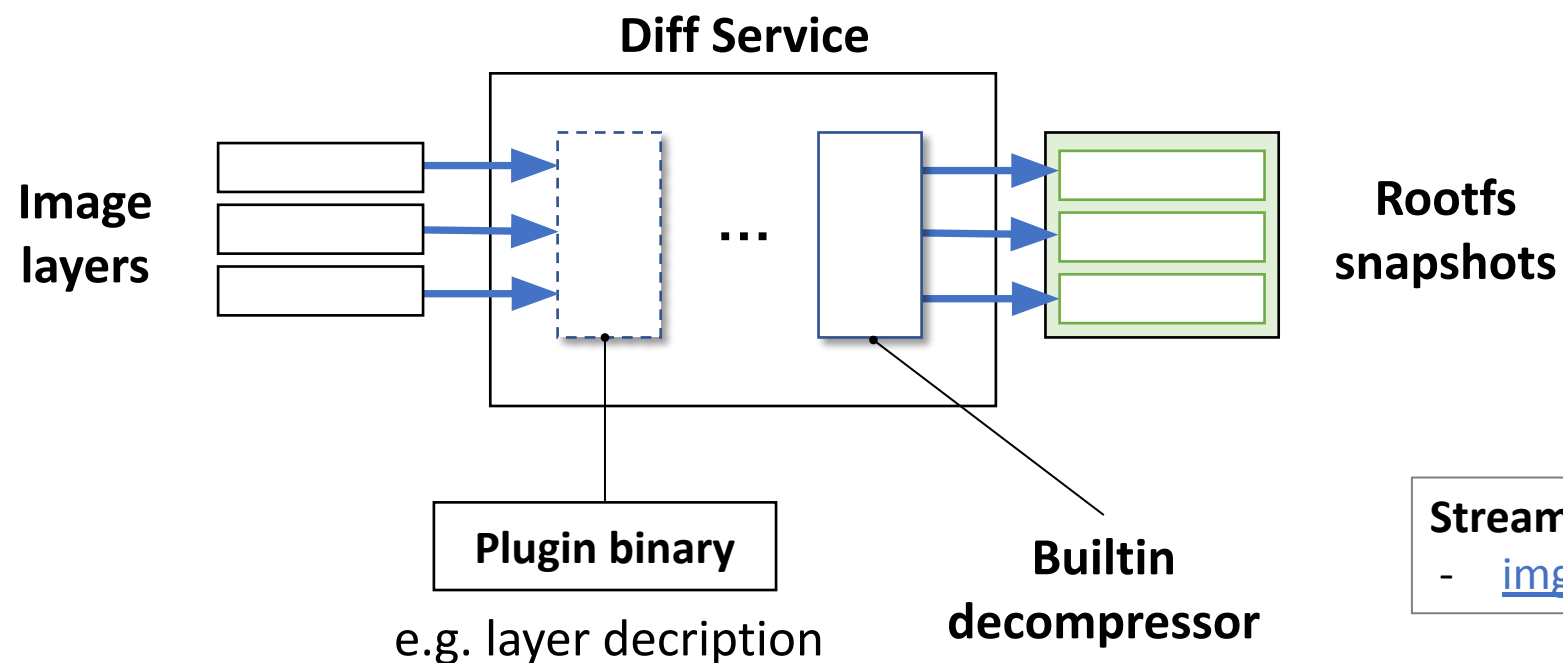


CloudNativeCon

Europe 2021

*Virtual*

- Containerd can handle arbitrary image layers, not limited to OCI standards
  - gzip, zstd, encrypted layers...
- Stream Processor plugin converts arbitrary media type to another (e.g. OCI standard types)
- Separated binary can plug into containerd, without re-compilation



**Stream Processor in community**  
- [imgcrypt](#) for encrypted images

## Extension example 3: Integrating low-level runtimes



KubeCon

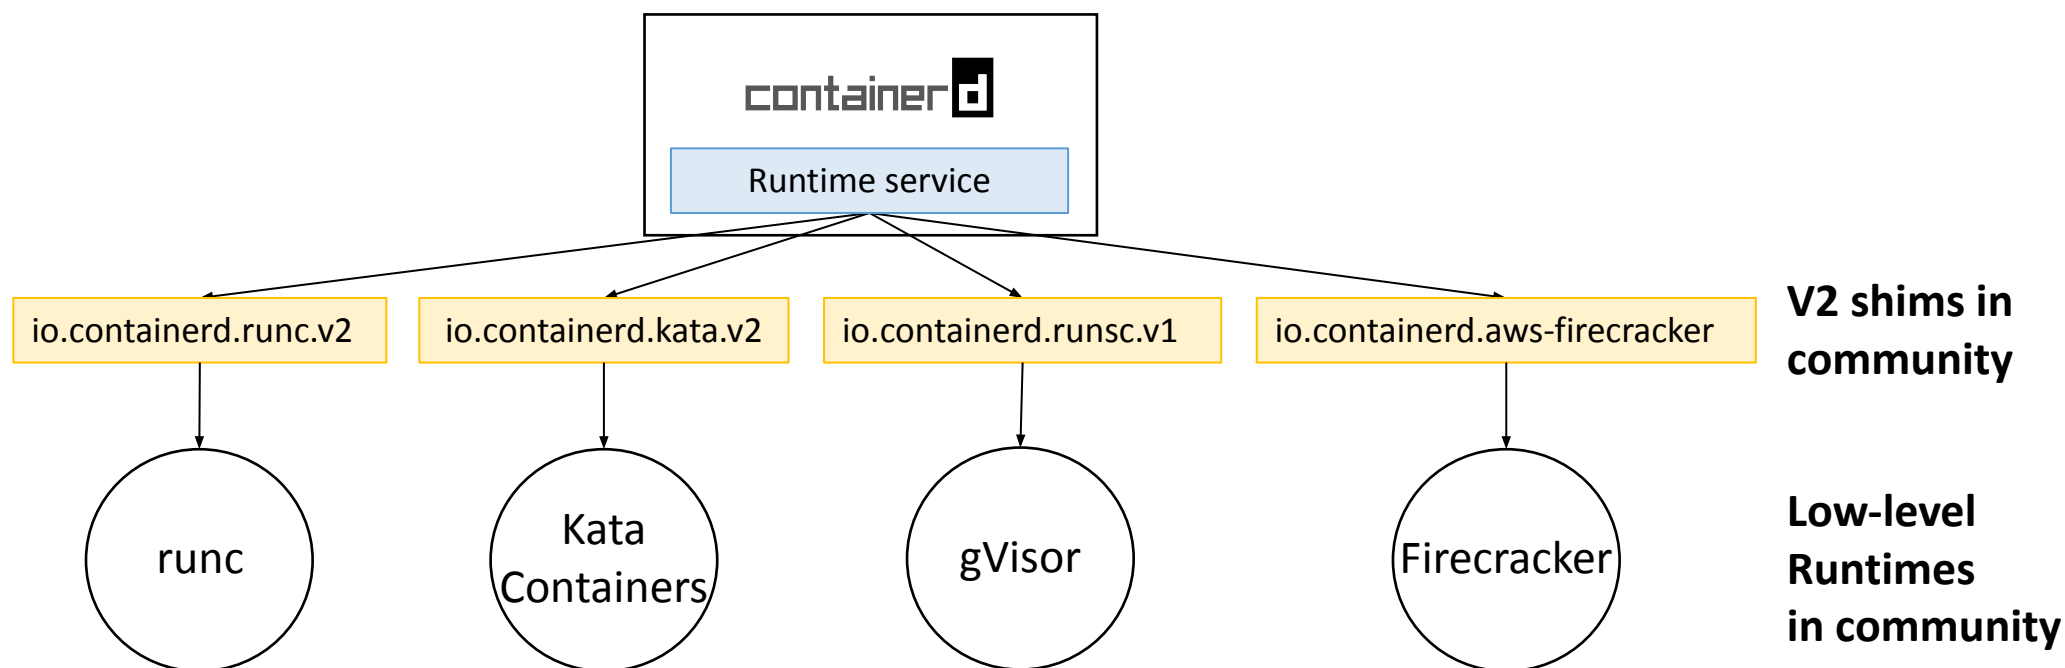


CloudNativeCon

Europe 2021

*Virtual*

- V2 Shim per low-level runtime
- Both of OCI (e.g. runc) and Non-OCI (e.g. Firecracker) runtime can integrate to containerd
- Binary naming convention: io.containerd.runc.v2 -> containerd-shim-runc-v2
- Pluggable logging destination
  - fifo(Linux), npipe(Windows), external binary(Linux, Windows), file(Linux, Windows)





KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# Implementing your own containerd client

Akihiro Suda, NTT Corporation

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

Two APIs are available

	containerd API	CRI API
<b>Consumers</b>	Docker/Moby, BuildKit, faasd, nerdctl...	Kubernetes
<b>Paradigm</b>	Container-oriented	Pod-oriented
<b>Flexibility</b>	Good	Bad
<b>Simplicity</b>	Bad	Good
<b>Transportation</b>	gRPC over UNIX socket	gRPC over UNIX socket

containerd API is recommended for most use cases, but CRI API might be easier to get started

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- Both containerd API and CRI API use gRPC
- In theory you could use any language for your own client
- But containerd API depends on “smart client” written in Go, especially for pulling images
- So, currently, Go is the best language for Native API
- Contribution is wanted for other languages



# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

Example: <https://containerd.io/docs/getting-started/>

```
client, err := containerd.New("/run/containerd/containerd.sock")
if err != nil {
    return err
}
defer client.Close()

// create a new context with an "example" namespace
ctx := namespaces.WithNamespace(context.Background(), "example")

// pull the redis image from DockerHub
image, err := client.Pull(ctx, "docker.io/library/redis:alpine", containerd.WithPullUnpack)
```

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

Example: <https://containerd.io/docs/getting-started/>

```
container, err := client.NewContainer(  
    ctx,  
    "redis-server",  
    containerd.WithImage(image),  
    containerd.WithNewSnapshot("redis-server-snapshot", image),  
    containerd.WithNewSpec(oci.WithImageConfig(image)),  
)
```

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

Example: <https://containerd.io/docs/getting-started/>

```
container, err := client.NewContainer(  
    ctx,  
    "redis-server",  
    containerd.WithImage(image),  
    containerd.WithNewSnapshot("redis-server-snapshot", image),  
    containerd.WithNewSpec(oci.WithImageConfig(image)),  
)
```

You will add WithXXX options here:

- oci.WithProcessArgs
- oci.WithMounts
- oci.WithMemoryLimit
- seccomp.WithProfile
- ...

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

In addition to the client, you will also want to implement OCI hooks and logger binary

- **OCI Hooks:** custom commands called on creation and deletion of containers
  - e.g., for setting up and tearing down CNI bridge and portmap
  - Optional, but necessary if you want your containers to be restarted automatically on host reboot
  - Example: <https://github.com/containerd/nerdctl/blob/v0.7.2/run.go#L629-L663>
- **Logger Binary:** custom command for handling container logs
  - e.g., store as a local file, transfer to fluentd, ...
  - Example: <https://github.com/containerd/nerdctl/blob/v0.7.2/run.go#L618-L627>

# Implementing your own containerd client



KubeCon



CloudNativeCon

Europe 2021

Virtual

Full example: **nerdctl**

<https://github.com/containerd/nerdctl>

Spun out from `ctr` tool with more practical features:

- Automatic restarting
- Port forwarding
- Logging
- Rootless
- Stargz
- OCIncrypt
- ...

You may copy the code as the “**starter pack**” to create your own client :)

[\[Download\]](#) [\[Command reference\]](#) [\[Additional documents\]](#)

## nerdctl: Docker-compatible CLI for containerd

nerdctl is a Docker-compatible CLI for [containerd](#).

- ✓ Same UI/UX as [docker](#)
- ✓ Supports [rootless mode](#)
- ✓ Supports [lazy-pulling \(Stargz\)](#)
- ✓ Supports [encrypted images \(ocrypt\)](#)

nerdctl is a **non-core** sub-project of containerd.

### Examples

#### Basic usage

To run a container with the default CNI network (10.4.0.0/24):

```
# nerdctl run -it --rm alpine
```

To build an image using BuildKit:

```
# nerdctl build -t foo .  
# nerdctl run -it --rm foo
```



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

# containerd 1.5 updates and future plan

Akihiro Suda, NTT Corporation

# containerd 1.5 updates (April)



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- Support **zstd** as an image compression algorithm
  - Faster than gzip
  - <https://facebook.github.io/zstd/>
- Support **NRI**: Node Resource Interface
  - Akin to CNI, but for managing resources, e.g., cgroup
  - <https://github.com/containerd/nri>
- Enable **OClcrypt** decryption by default
  - Supported since 1.3, but it was not enabled by default
  - <https://github.com/containerd/ocicrypt> <https://github.com/containerd/imgcrypt>
- **nerdctl** (contaiNERD ctl) joined containerd, as a non-core subproject
  - Docker-compatible CLI but with stargz and ocicrypt
  - <https://github.com/containerd/nerdctl>

# containerd 1.5 updates (April)



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- The CRI plugin repo ([github.com/containerd/cri](https://github.com/containerd/cri)) is now merged into the main repo ([github.com/containerd/containerd](https://github.com/containerd/containerd))
  - No visible change to users, but significantly simplifies contribution process
- Client library is now available as a Go module



# Future plan



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- Filesystem quota ([#759](#))
- CRI support for user namespaces ([KEP #2101](#))
  - Run Kubernetes pods as a user that is different from the daemon user
  - Akin to “Rootless Containers”, but different (and does not conflict, either)
- Chown-less user namespaces ([#4734](#))
  - Requires idmapped mounts, introduced in kernel 5.12
- Pause-less pod sandboxes ([#4131](#))
- More documentation (help wanted! 🙏)

# Third party plugin updates



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- Nydus Snapshotter <https://github.com/dragonflyoss/image-service>
  - Similar to Stargz Snapshotter but with a different image format
- OverlayBD Snapshotter <https://github.com/alibaba/accelerated-container-image>
  - Boot containers from iSCSI
- runu <https://github.com/ukontainer/runu>
  - Linux containers on macOS, using LKL (Linux Kernel Library)
- runj <https://github.com/samuelkarp/runj>
  - FreeBSD containers

# Recap



KubeCon



CloudNativeCon

Europe 2021

*Virtual*

- The de facto standard runtime for Kubernetes, but not only for Kubernetes
- Extensible with plugins
  - Runtime plugins, e.g., gVisor, Kata
  - Snapshotter plugins, e.g., Stargz Snapshotter
  - Stream processor plugins, e.g., OCIcrypt
  - Logging binary plugins, e.g., json-file
  - ...
- New subproject: nerdctl (<https://github.com/containerd/nerdctl>)
  - Like `docker` but with full features of containerd
  - Like `ctr` but with full user experience of `docker`
  - `nerdctl run -d -p 80:80 --restart=always nginx`



KubeCon



CloudNativeCon

Europe 2021

*Virtual*



*Forward Together »*