



KubeCon



CloudNativeCon

Europe 2021

Virtual



Forward Together »

Lessons Learned Deploying Traditional Web Applications on Top of Kubernetes

Marcos Bjoerkelund (He/Him)



KubeCon



CloudNativeCon

Europe 2021

Virtual



Agenda

Introduction

General Workflow

From Build to Deploy

Tips and Good Practices

To Create Better Container Images

Challenges You May Find

And Ways to Solve Them

An Example

Deploying WordPress on Kubernetes



KubeCon



CloudNativeCon

Europe 2021

Virtual

Introduction



KubeCon



CloudNativeCon

Europe 2021

Virtual

Bitnami is a Catalog of Free Open Source Software

10+ years building and maintaining software packages

Over 180 applications, runtimes, frameworks and more, including:



Any Environment



Local



Cloud



Data Center

Any Format



Virtual
Machines



Containers



Deployment
Templates

Any Platform



kubernetes



vmware



aws



Azure



Google



IBM Cloud



KubeCon



CloudNativeCon

Europe 2021

Virtual

Traditional Web Applications




Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data, lack of external input methods (i.e. env vars)



Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual


Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data, lack of external input methods (i.e. env vars)



Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual


Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data and config, lack of external input methods (i.e. env vars)



Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual


Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data, lack of external input methods (i.e. env vars)



Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual


Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data, lack of external input methods (i.e. env vars)




Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual


Traditional Web Applications

Do not follow key requirements for cloud-native apps (based on <https://12factor.net>):




State & Processes

Processes are stateless and share-nothing



Config mgmt.

No separation with code/data, lack of external input methods (i.e. env vars)



Disposability

Can be started/stopped at any moment, with minimal startup time



Concurrency

Application must be able to scale horizontally



Dependencies

Reproducible environments



KubeCon



CloudNativeCon

Europe 2021

Virtual

Examples

github.com/bitnami-labs/deploy-web-apps-on-kubernetes



KubeCon



CloudNativeCon

Europe 2021

Virtual[™]

General Workflow

From Build to Deploy



KubeCon



CloudNativeCon

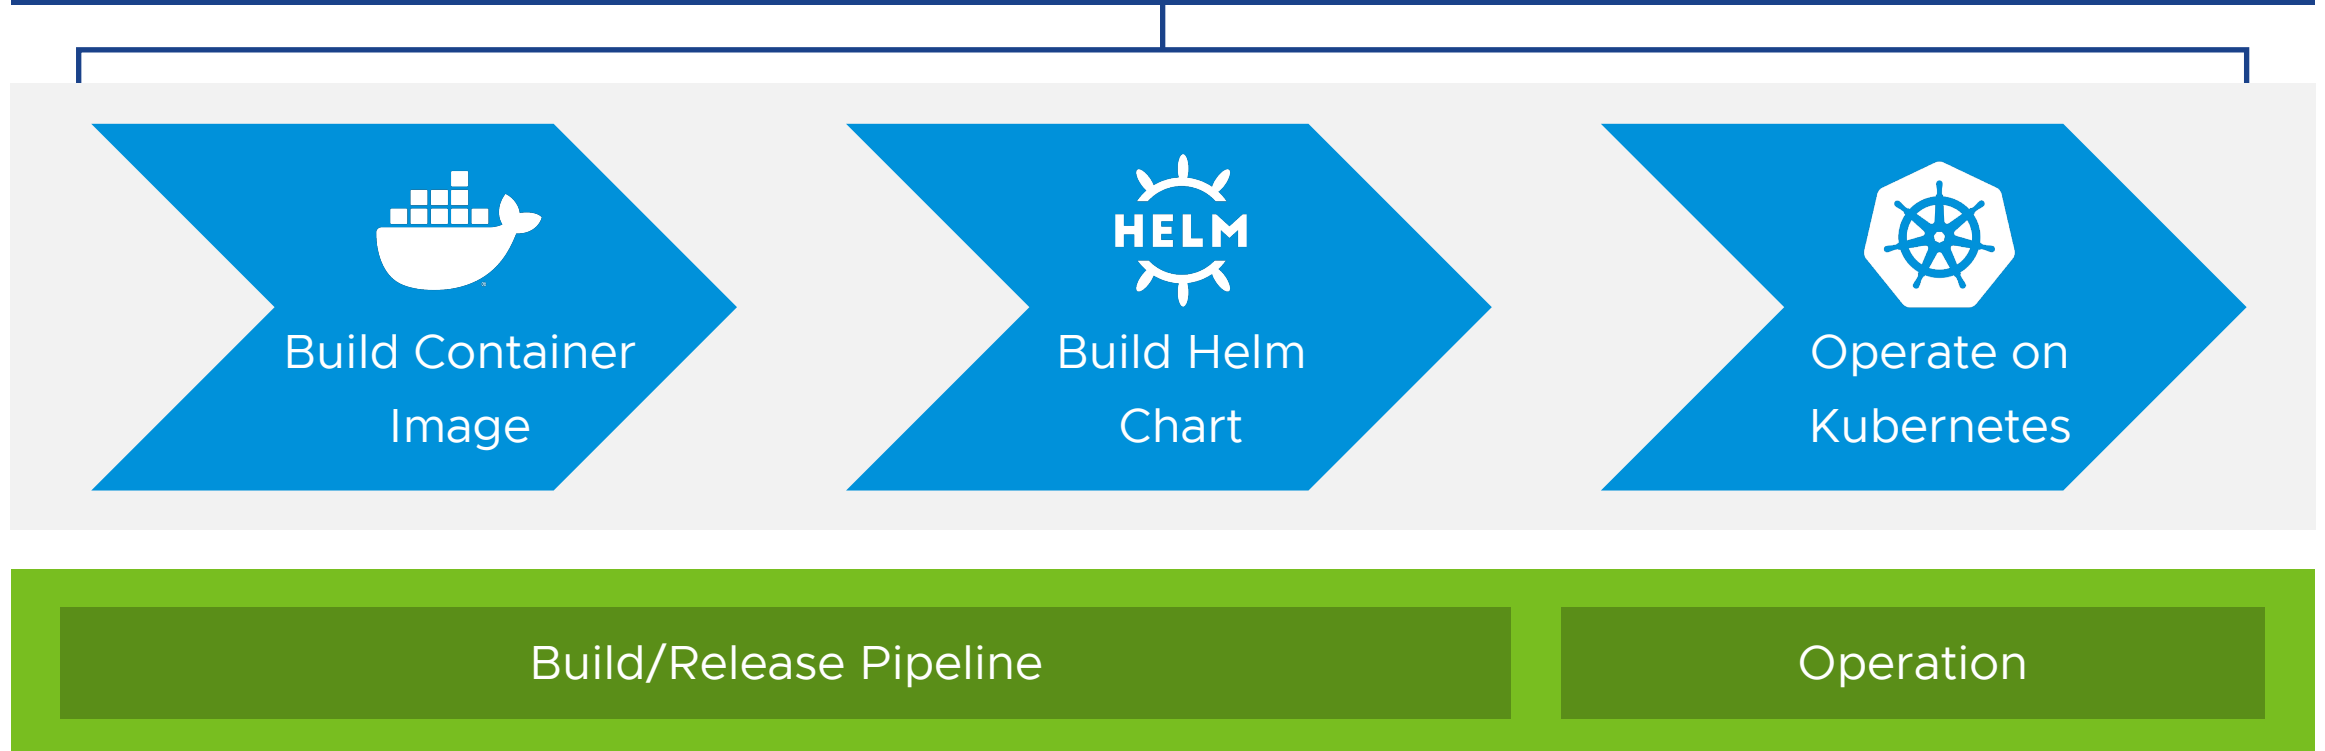
Europe 2021

Virtual

General Workflow for a Traditional Web Application

What processes and steps are required to deploy on Kubernetes

Workflow to deploy application on Kubernetes



KubeCon



CloudNativeCon

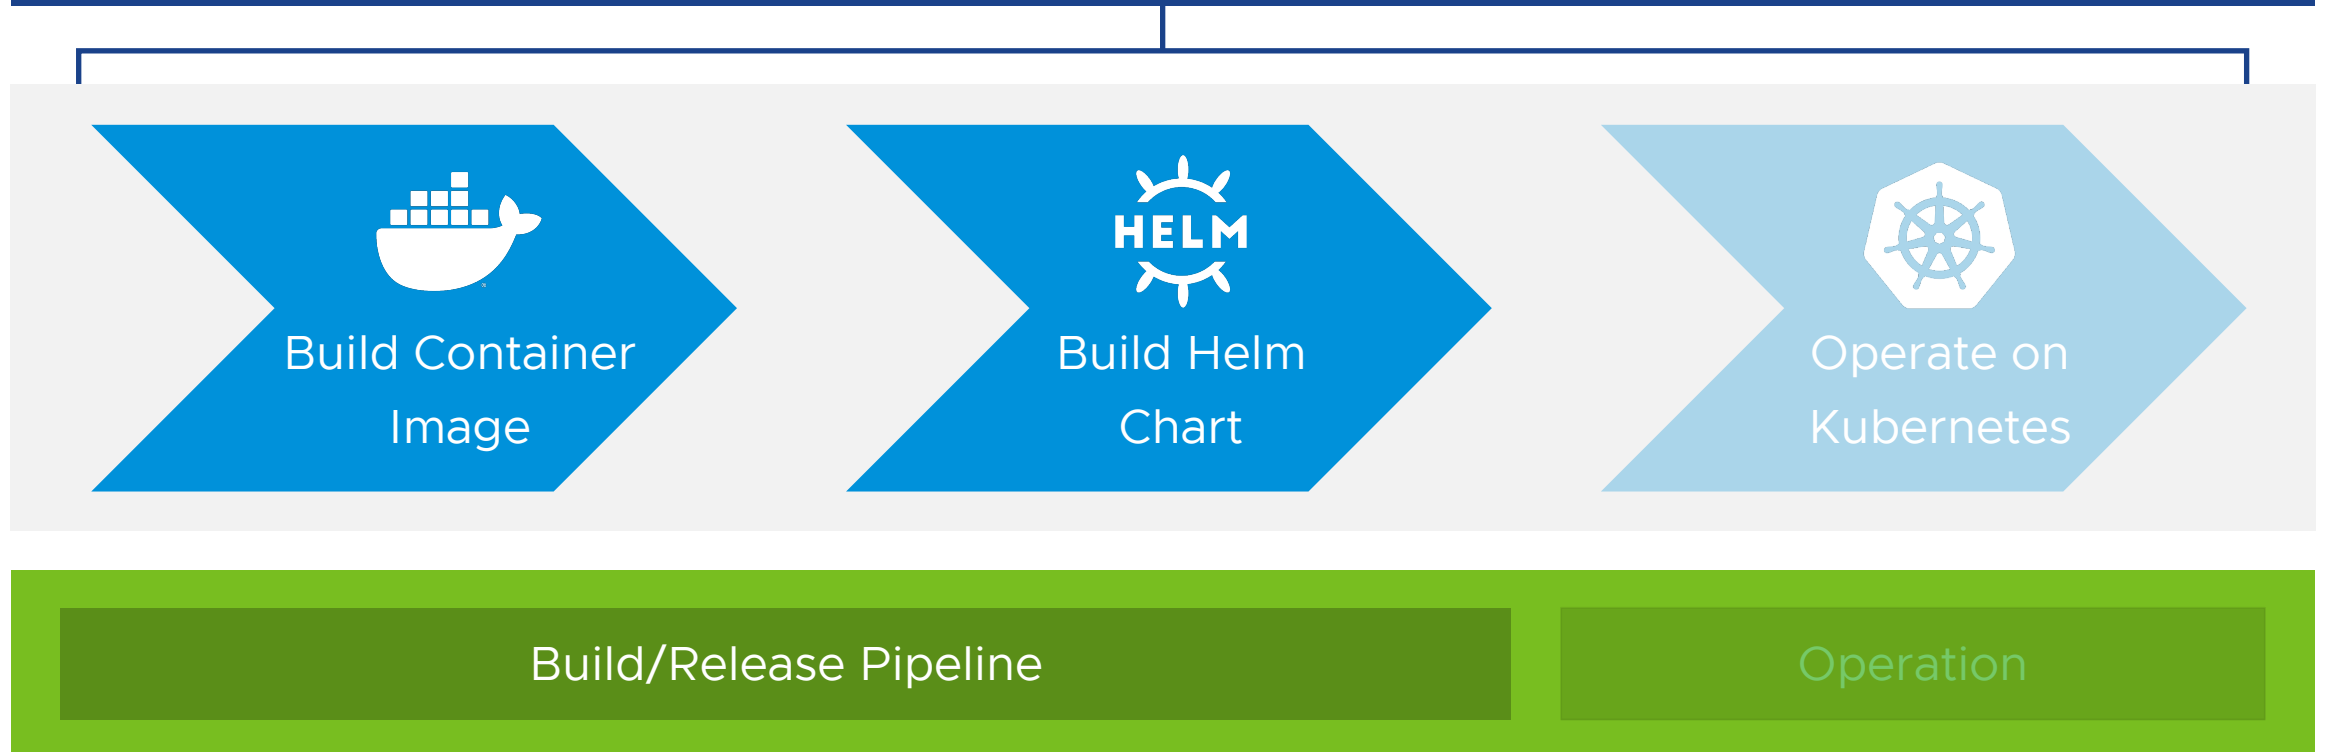
Europe 2021

Virtual

General Workflow for a Traditional Web Application

What processes and steps are required to deploy on Kubernetes

Workflow to deploy application on Kubernetes



KubeCon



CloudNativeCon

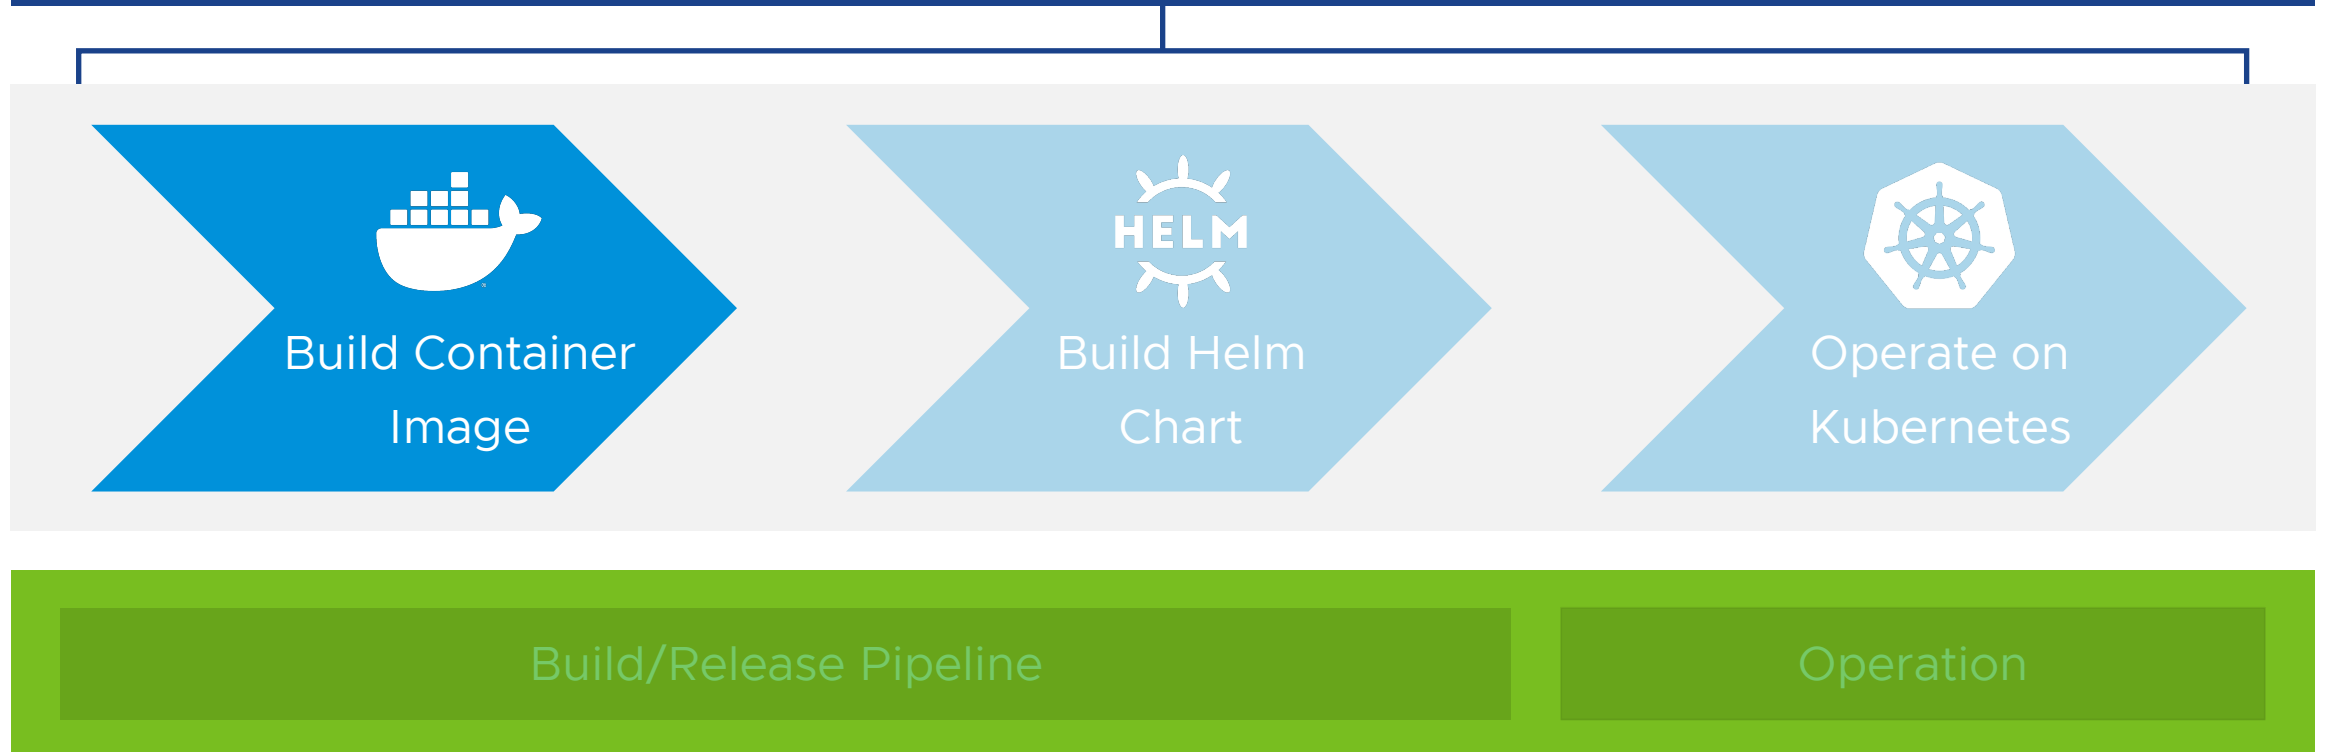
Europe 2021

Virtual

General Workflow for a Traditional Web Application

What processes and steps are required to deploy on Kubernetes

Workflow to deploy application on Kubernetes



KubeCon



CloudNativeCon

Europe 2021

Virtual

Build Container Image

Container images overview

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/1-container-image

- OCI compliant images (Docker, Podman)
- Contexts: Build-time & runtime
- Build-time (Dockerfile):
 - List of steps to build/run an image
 - Based on commands and layers
 - Limited input support (ARG)
- Runtime (entrypoint)
 - Logic to install and run the application
 - Inputs via env. vars (or mounted files)

Example: Dockerfile

<code>FROM</code> debian:buster	◆	Base image
<code>RUN</code> apt-get update && \	◆	Runtime
apt-get install -y python		deps
<code>COPY</code> html /app	◆	Install app
<code>WORKDIR</code> /app		
<code>EXPOSE</code> 8080	◆	Runtime
<code>USER</code> www-data		config
<code>CMD</code> python -m SimpleHTTPServer 8080		

Example: Build image from Dockerfile

```
$ docker build -t image:tag .
```



KubeCon



CloudNativeCon

Europe 2021

Virtual

Build Container Image

Dockerfile design

Must-have

Install dependencies and system package

Copy and prepare application files

Establish **image security**

Good to have

Focus on **minimal image size**

Design with **runtime stability** in mind, e.g.:

- CPU- and memory-intensive steps at build
- Pre-download files from the Internet

Runtime configuration defaults



KubeCon



CloudNativeCon

Europe 2021

Virtual

Build Container Image

Entrypoint design

Must-have

Focus on **minimal deployment time**

Validation of inputs (env. vars / filemounts)

Persistence of application data

Start web application

Good to have

Configuration logic (dependent on inputs)

Wait for external services

- Avoid potential app errors and warnings

Pre-/post-installation features, e.g.:

- Further extend functionality of the app
- Automatic installation in the database



KubeCon



CloudNativeCon

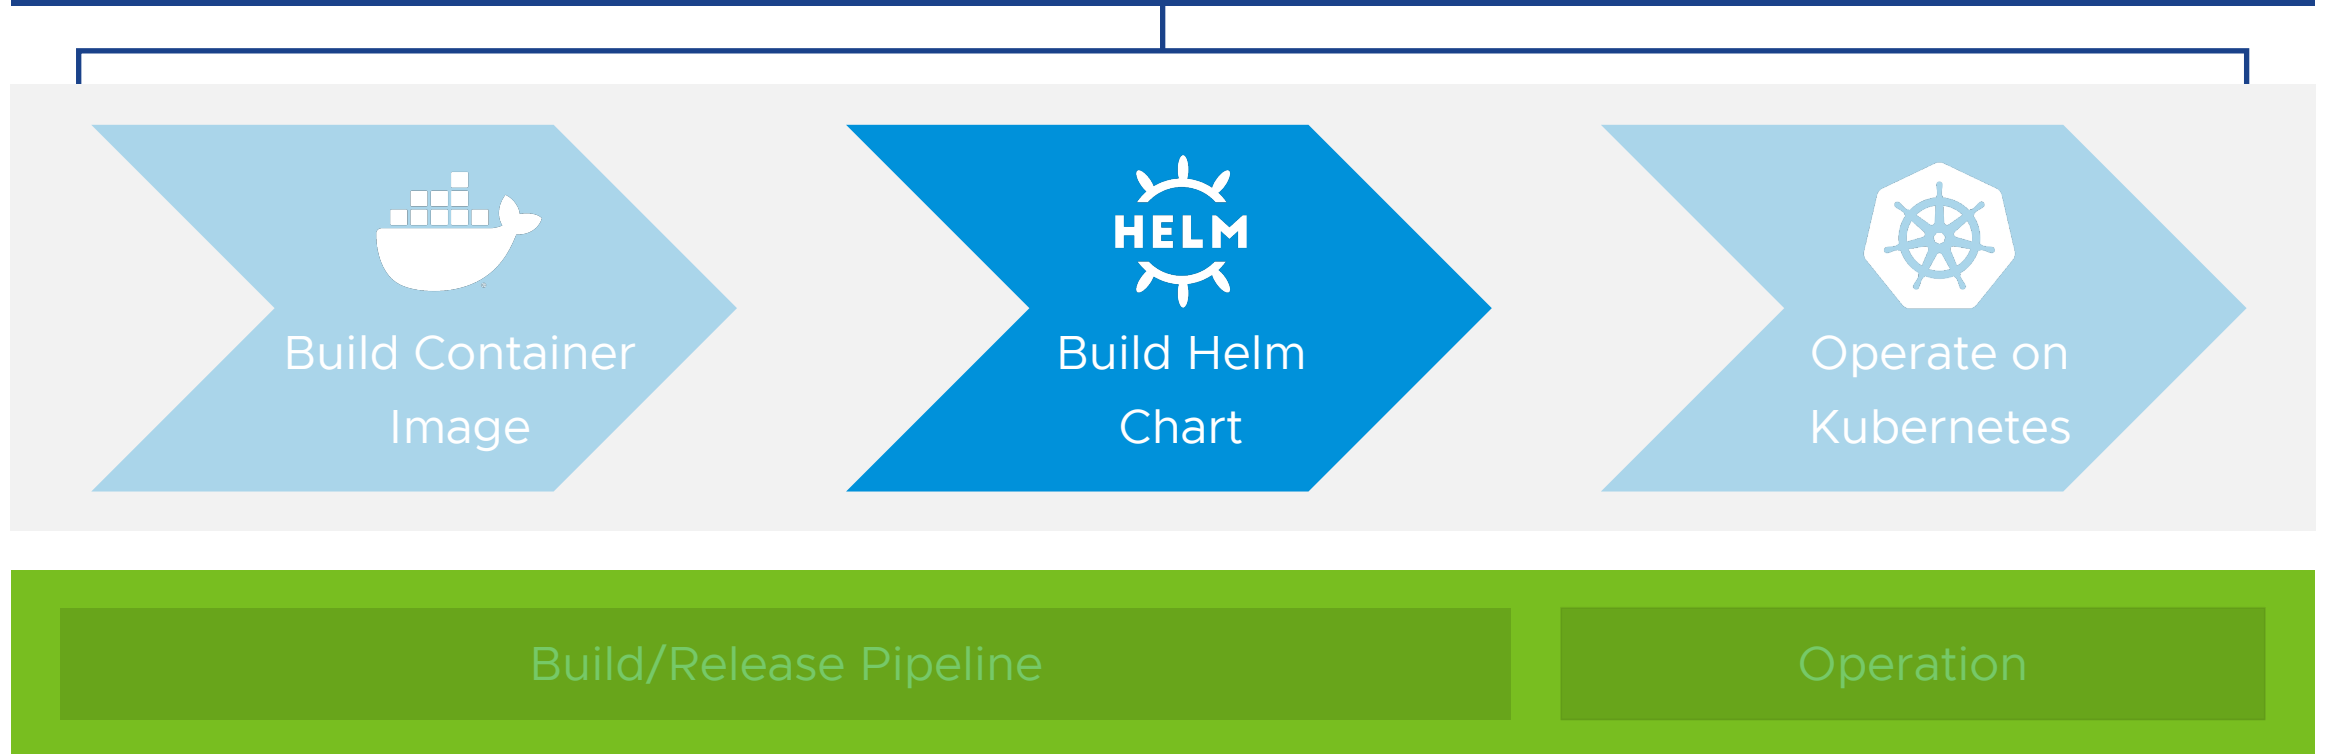
Europe 2021

Virtual
20

General Workflow for a Traditional Web Application

What processes and steps are required to deploy on Kubernetes

Workflow to deploy application on Kubernetes



KubeCon



CloudNativeCon

Europe 2021

Virtual²¹

Build Helm Chart

Helm charts overview

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/2-helm-chart

Introduction to [Helm](#):

- Most popular for deploying web apps on K8s
- Defines how the application should operate
- Inputs via values.yaml or CLI (--set) ¹
- Template of Kubernetes resources in YAML ²
- Sub-charts? ³

To deploy a Helm chart on Kubernetes:

```
$ helm install my-release myrepository/wordpress \
  --set option1=value1[,option2=value2]
```

Example: File tree of a chart

```
wordpress/
  Chart.yaml
  LICENSE
  README.md
  values.yaml 1
  values.schema.json
  templates/ 2
  templates/NOTES.txt
  charts/ 3
  crds/
```



KubeCon



CloudNativeCon

Europe 2021

Virtual

Build Helm Chart

Helm chart design

Essentials

Deployment of for app containers

Ingress with **TLS/SSL**

Persistent volume claims with app data

Secrets with credentials

...

Good to have

Database or external service,
either external or sub-chart

Metrics support (service, container, ...)

Cron jobs for temporary jobs

Custom resources, init scripts...

...



KubeCon



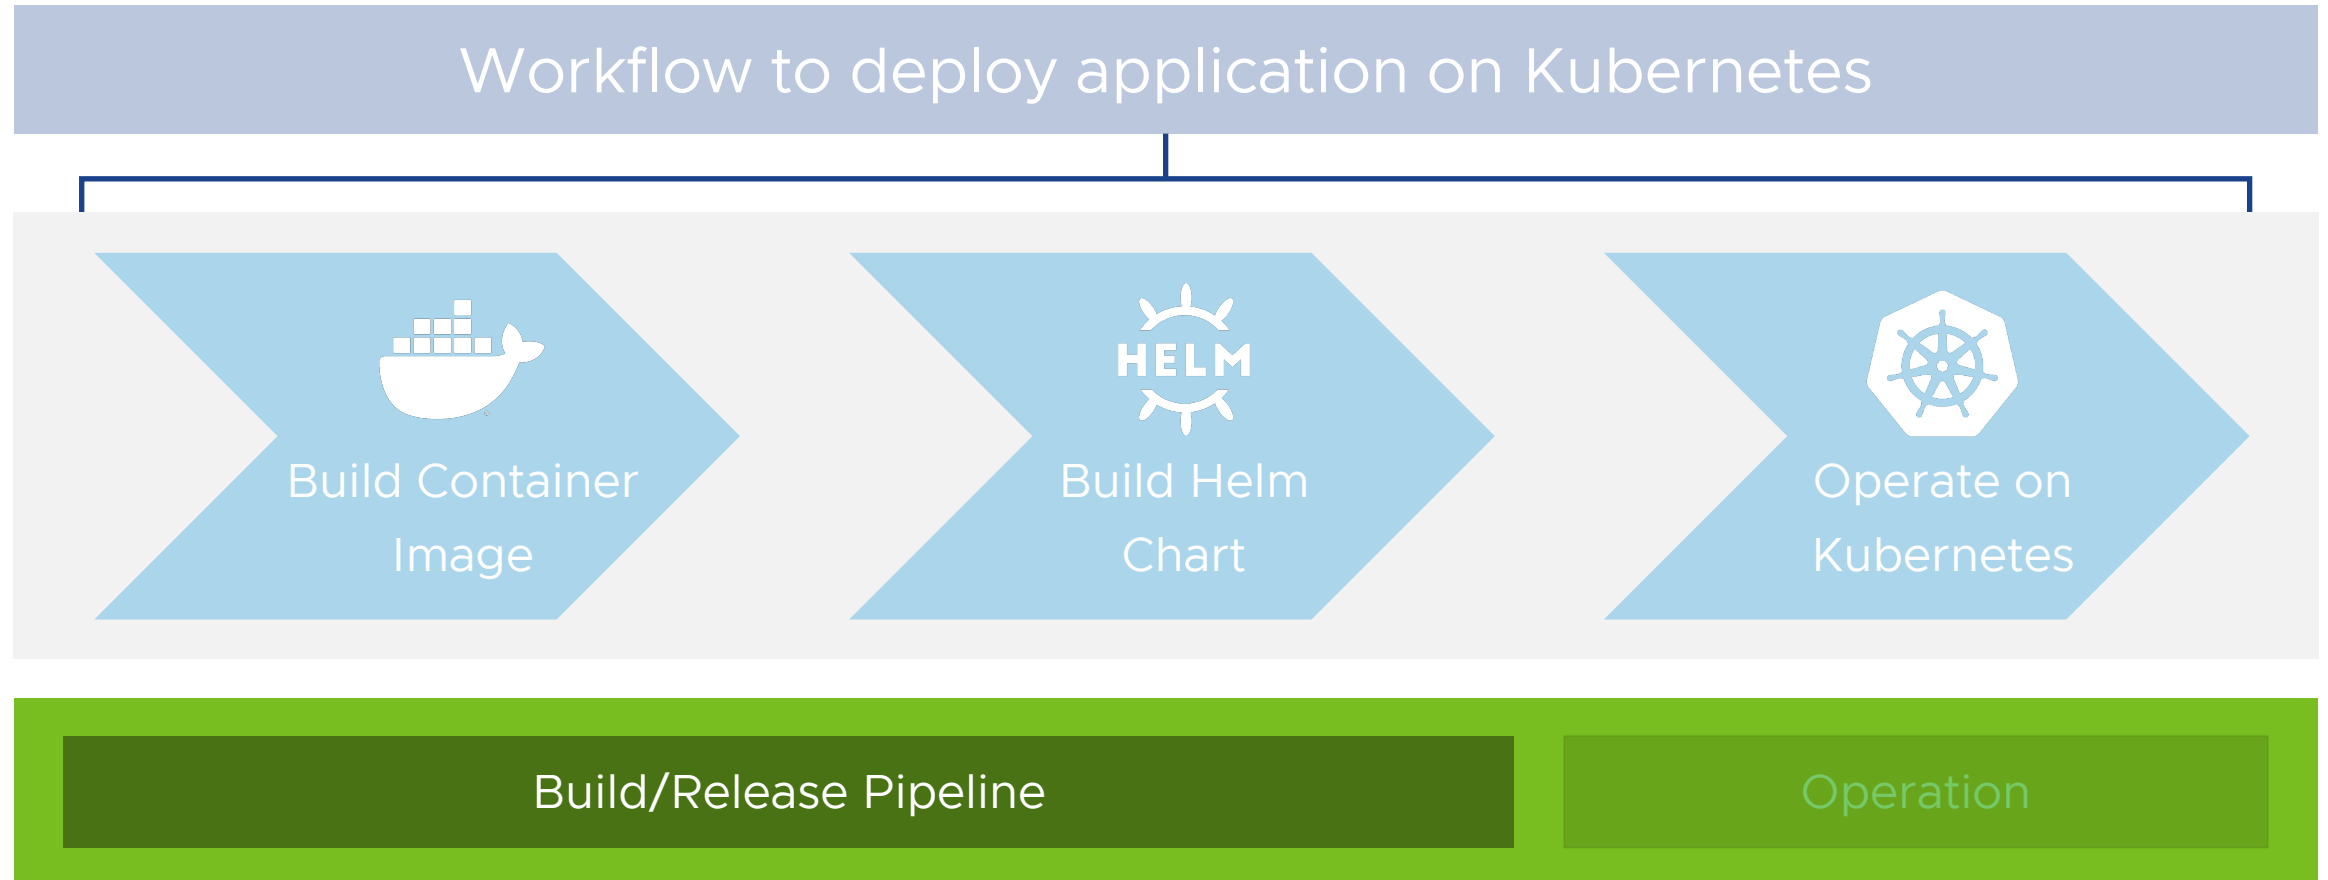
CloudNativeCon

Europe 2021

Virtual

General Workflow for a Traditional Web Application

What processes and steps are required to deploy on Kubernetes



KubeCon



CloudNativeCon

Europe 2021

Virtual

Build and Release Pipeline

What steps to perform on a build and release pipeline

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/3-pipeline

Basic steps:

- Build container image (“docker build”)
- Build Helm chart (“helm package”)
- Testing container image & chart
- Release: Publish images and charts

Multiple free (*) and simple CI/CD options available:

- [GitHub Actions](#)
- [Travis CI](#)
- [Jenkins](#) / [Jenkins X](#)
- [GitLab CI](#)
- [CircleCI](#)
- And many more

(*): There may be limits to the free tier, except from self-hosted options (like Jenkins or GitLab CE).

Bitnami’s pipeline:



See: [“App Testing at Scale: How Bitnami Tests Thousands of Releases Per Month” by Juan Jose Martos Castro](#)

(KubeCon/CloudNativeCon North America 2020)



KubeCon



CloudNativeCon

Europe 2021

Virtual

Tips and Good Practices

To Create Better Container Images and Helm Charts



KubeCon



CloudNativeCon

Europe 2021

Virtual

Security

Non-root users

Non-root containers protect from **privilege escalations** and **access to non-writable files**



Existing non-root username (known UID)

- Incompatible with platforms that require non-root or arbitrary UID (i.e. OpenShift)
- Writable files: Explicit file/group ownership



Existing non-root username (arbitrary UID)

- Works in all environments
- Writable files: Explicit group ownership
- **Insecure:** Requires writable /etc/passwd

[Link to OpenShift docs](#)



Arbitrary UID (no username)

- Works in all environments
- Incompatible with some apps (workaround: nss_wrapper)
- Ugly prompts (“I have no name!”)
- Writable files: g+rwX permissions (root owner)



KubeCon



CloudNativeCon

Europe 2021

Virtual

Security

Non-root users

Non-root containers protect from **privilege escalations** and **access to non-writable files**



Existing non-root username (known UID)

- Incompatible with platforms that require non-root or arbitrary UID (i.e. OpenShift)
- Writable files: Explicit file/group ownership



Existing non-root username (arbitrary UID)

- Works in all environments
- Writable files: Explicit group ownership
- **Insecure:** Requires writable /etc/passwd

[Link to OpenShift docs](#)



Arbitrary UID (no username)

- Works in all environments
- Incompatible with some apps (workaround: nss_wrapper)
- Ugly prompts (“I have no name!”)
- Writable files: g+rwX permissions (root owner)



KubeCon



CloudNativeCon

Europe 2021

Virtual

Security

Permissions

- Permissions should be set at build time, not at runtime:
 - **Limited** ability to set permissions at runtime by non-root users
 - **Slow**, affecting deployment time (at build time we have all the time in the world)
- Limit list of writable files to the bare minimum (i.e. only data directories).
- Advanced topics: [Docker Security Cheat Sheet \(OWASP\)](#)
 - Read-only file systems
 - Use Linux Security Modules: AppArmor, seccomp, SELinux
 - ...



KubeCon



CloudNativeCon

Europe 2021

Virtual

Optimize Container Image Size

How to reduce your container image size

The **layer format** of containers images may have a **huge impact** in their size

Motivation: Reduce bandwidth usage and deployment time (*).

- Multi-stage image (especially for apps w/ build-time deps or must be compiled)
- Strategic Dockerfile design: Group commands + Remove unused files in each step + Location
- Docker experimental feature: `--squash` – Beware of its [limitations](#)

Useful tool: [wagoodman/dive](#)

```
| Image Details |  
Image name: bitnami/wordpress:latest  
Total Image size: 1.6 GB  
Potential wasted space: 1.1 GB  
Image efficiency score: 54 %
```



```
| Image Details |  
Image name: bitnami/wordpress:optimized  
Total Image size: 602 MB  
Potential wasted space: 4.0 MB  
Image efficiency score: 99 %
```



KubeCon



CloudNativeCon

Europe 2021

Virtual

Good practices

Create container images following good practices guidelines

- [Best Practices Writing a Dockerfile](#)
- 1 process per container!
- Run heavy and time-consuming tasks needed to bootstrap the application at build-time
 - E.g. build assets, install npm dependencies
- Reproducibility: Beware of network issues, distro package manager changes, etc.
- Maintenance: Periodic re-build of the image
 - Minimize response time in case of a security incident of an application dependency
 - Detect new build-time issues early before a re-build is urgent
- Use simple tools for initializing the app (e.g. Bash): Visibility, easy to debug, little memory usage
- Other: Multi-arch, distro election depending on your needs...



KubeCon



CloudNativeCon

Europe 2021

Virtual
31

Good practices

Create Helm charts following good practices guidelines

- [Best Practices for Creating Production-Ready Helm Charts](#)
- [Bitnami Helm chart template following good practices](#) (charts/template)
- Deploy non-root containers by default.
- Extensibility and flexibility. For example:
 - Extend the app configuration.
 - Components and features to enable/disable via chart inputs.
 - Custom init containers, volumes, environment variables, annotations, resource limits, etc.
 - Custom additional Kubernetes resources to deploy.
- Integrate with logging and monitoring platforms.
- Diagnostics: Make it easier to debug deployments in case of a failure.



KubeCon



CloudNativeCon

Europe 2021

Virtual

Challenges You May Find

And Ways to Solve Them



KubeCon



CloudNativeCon

Europe 2021

Virtual

Persistence of Application Data

What application files to persist between runs

Only persist data files that must be writable at runtime. Separation of data & code.



App separates code and data

Persist **certain** application files (data)

Not supported for all applications

Faster as less data must be populated.

Automatic app update process support



No separation of code and data

Persist **all** application files

Supported for all applications

Slower as it requires volume population.

Limited maintenance & upgrade support



KubeCon



CloudNativeCon

Europe 2021

Virtual
IN

Application Updates

How to implement support for automatic application updates

App must separate code and data, or provide a UI interface for upgrading.



Automatic updates support

Backup DB before upgrading! Scale down!

App separates **code and data**

- Update version: Swap image version (tag).

App provides **web UI update tool**

- Force insecure permissions.
- No rollback in case of errors.



Requires manual updates

May be your **only choice**:

- Unfeasible to automate upgrade script.
- Code and data separation is not supported.



KubeCon



CloudNativeCon

Europe 2021

Virtual

Multiple Application Services

How to support running multiple application services

Avoid running > 1 service per container

- Problem: Each service require shared data with the app (i.e. workers):
 - Does the app support external services for cache/data? (i.e. Redis)
 - Alternative: Synchronized / cached FS (i.e. NFS) if data needs to be available for all services
- Problem: Frontend and backend servers running the same code:
 - CDN for storing non-static assets (frontend), route all requests to backend
 - Shared / cached file systems (see above)
 - Varnish: Serve everything through backend, cache repeatable requests to reduce app load
- Other cases:
 - Periodic jobs (i.e. Crontab) to use Kubernetes CronJob resource with app volumes.
 - Init containers for pre-install steps (i.e. further customize your application), aside from container logic.



KubeCon



CloudNativeCon

Europe 2021

Virtual[™]

Horizontal Scaling

How to design a Helm chart to support scaling your application up/down

- Ideally: Rely on external services, no local data. For example:
 - Web server to run the application scripts.
 - Server for temporary files (i.e. Redis or Memcached): cache, sessions.
 - CDN: host assets, user uploads.
 - External load balancer as main entrypoint for the application. Headless services.
- Reality:
 - Limited scale support but require local data: Workaround via synchronized FS (previous slide).
 - May be impossible: Cache server (Varnish). Optimize application configuration. Scale vertically.



KubeCon



CloudNativeCon

Europe 2021

Virtual

An Example

Deploying WordPress on Kubernetes



KubeCon



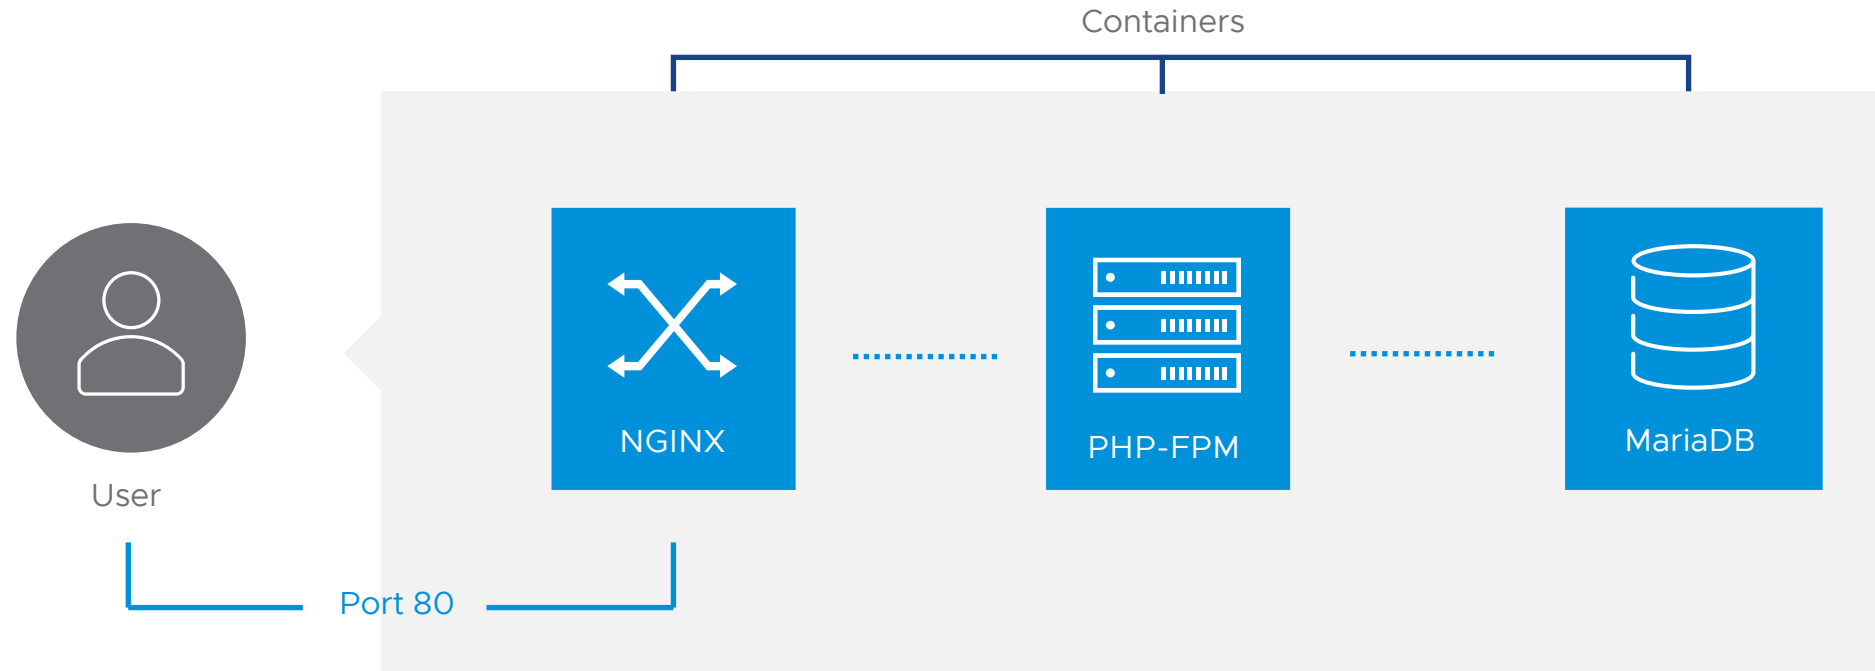
CloudNativeCon

Europe 2021

Virtual

Service Diagram

How the different services will interact with each other



KubeCon



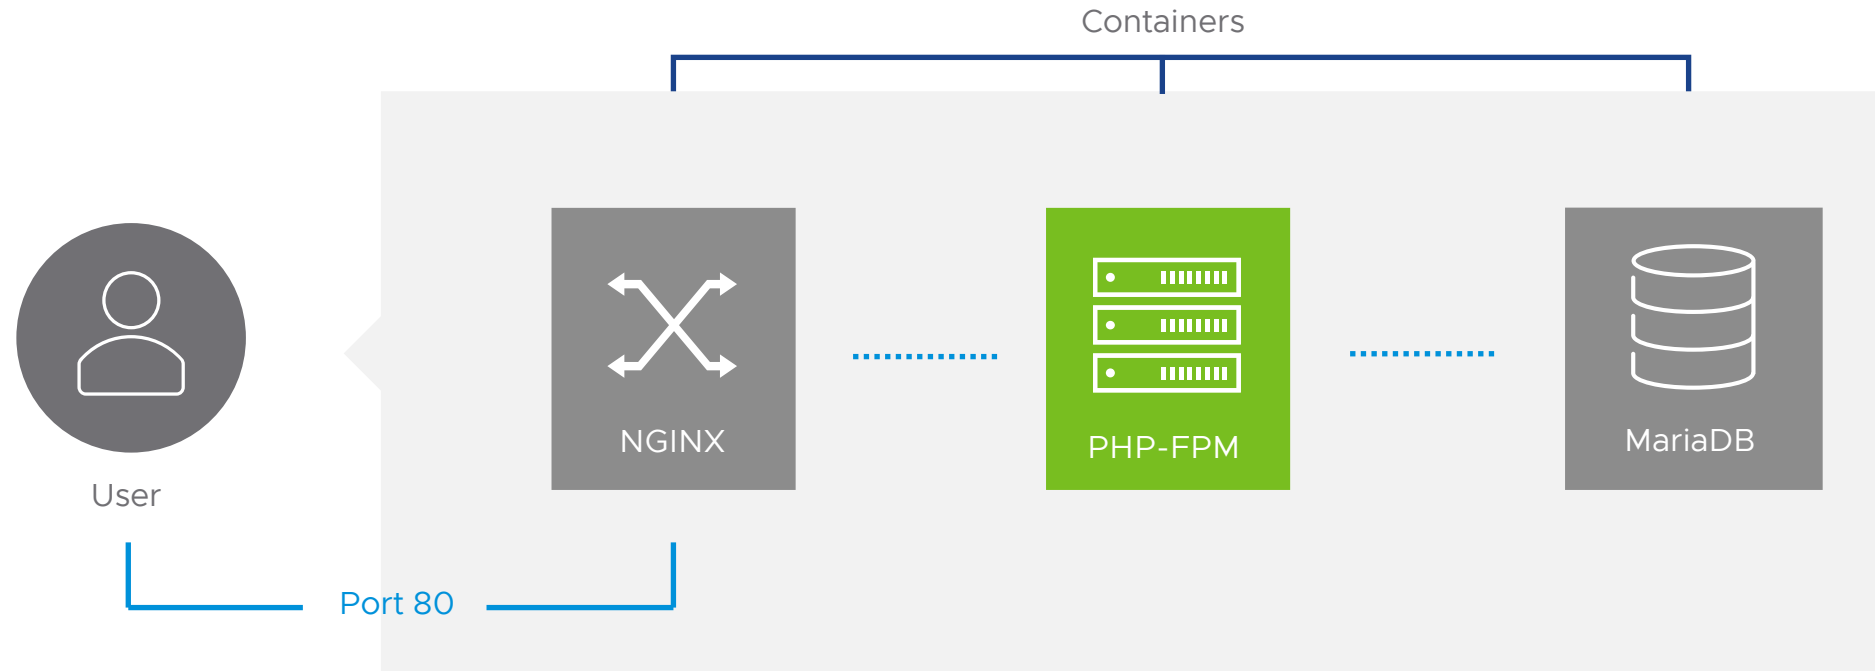
CloudNativeCon

Europe 2021

Virtual

WordPress Deployment Service Diagram

How the different services will interact with each other



KubeCon



CloudNativeCon

Europe 2021

Virtual

WordPress Dockerfile

How the Dockerfile will be structured

FROM bitnami/php-fpm:8.0 AS builder	◆	Multistage build (same image to save space)
ARG WORDPRESS_VERSION=5.7		
RUN apt-get update && apt-get install -y ca-certificates curl	◆	Install build dependencies (for cURL)
RUN curl -s https://wordpress.org/wordpress-\${WORDPRESS_VERSION}.tar.gz \	◆	Extract + permissions/ownership (security)
tar xz -C / --no-same-owner --no-same-permissions		
# RUN curl http://.../plugin.zip tar xz -C /wordpress/wp-content	◆	Build/customize WordPress installation
FROM bitnami/php-fpm:8.0	◆	Minimal image
RUN sed -Ei 's/;(clear_env)/\1/' /opt/bitnami/php/etc/php-fpm.d/www.conf	◆	Adapt base image to our needs
RUN apt-get update && apt-get install -y --no-install-recommends vim \	◆	Install runtime dependencies & clean cache
&& rm -r /var/lib/apt/lists /var/cache/apt/archives		
COPY rootfs /		
RUN mkdir /wp-content && chmod g+rwX /wp-content		
COPY --from=builder /wordpress /wordpress	◆	Install WordPress in final container
WORKDIR /wordpress		
ENTRYPOINT ["/entrypoint.sh"]	◆	Runtime configurations
CMD ["php-fpm", "-F"]		
VOLUME ["/wp-content"]	◆	Only persist "wp-content"
USER 1001	◆	Run as non-root user (security)
...		

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/1-container-image



KubeCon



CloudNativeCon

Europe 2021

Virtual

WordPress Entrypoint Script

How the different services will interact with each other

```
#!/bin/bash
Set -euo pipefail
. /functions.sh
if [[ "$1" = "php-fpm" ]]; then
    # Persistence of application data
    if is_dir_empty /wp-content; then
        # Validate inputs
        assert_not_empty WORDPRESS_DB_HOST
        ...
        assert_check
        info "Persisting WordPress content"
        cp -r /wordpress/wp-content/. /wp-content
    else
        info "Deploying WordPress with persisted content"
    fi
fi
info "Starting WordPress server"
exec "$@"
```

Annotations:

- Bash strict mode and load functions
- Determine if app is being initialized
- Validate inputs
- Persist application data files (1st run only)
- Already persisted data: Quick start (+ upgrade?)
- Start app server

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/1-container-image



KubeCon



CloudNativeCon

Europe 2021

Virtual

WordPress deployment on Kubernetes

Apache + mod_php example: [bitnami/wordpress](#) chart:

```
$ helm install my-wordpress bitnami/wordpress --set image.name=marcosbc/wordpress-modphp-example:5.7
```

- [Parameters](#)
- [MariaDB sub-chart](#), ability to specify an external DB
- Optional Memcached support

NGINX + PHP-FPM example:

github.com/bitnami-labs/deploy-web-apps-on-kubernetes/tree/main/2-helm-chart



KubeCon



CloudNativeCon

Europe 2021

Virtual



KubeCon



CloudNativeCon

Europe 2021

Thank You!

Virtual

